

EXTENDED OFFER | Flat 15% OFF SITEWIDE | Use Coupon - WHIZOFFER15



Search Courses



## Practice Test 1

<b>Attempt</b>	<b>Completed on</b>
5	Friday , 19 July 2019 , 07:29 AM
<b>Marks Obtained</b>	<b>Time Taken</b>
55 / 55	00 H 16 M 41 S
<b>Your score is</b>	<b>Result</b>
100%	Pass
<b>Mode</b>	
Exam	

## Domains wise Quiz Performance Report

<b>No</b>	1
<b>Domain</b>	Connect to and Consume Azure Services and Third Party Services
<b>Total Question</b>	16
<b>Correct</b>	16
<b>Incorrect</b>	0
<b>Unattempted</b>	0
<b>Marked for review</b>	0

<b>No</b>	2
<b>Domain</b>	Develop for Azure Storage
<b>Total Question</b>	11
<b>Correct</b>	11
<b>Incorrect</b>	0
<b>Unattempted</b>	0
<b>Marked for review</b>	0
<b>No</b>	3
<b>Domain</b>	Develop Azure Infrastructure as a Service Compute Solutions
<b>Total Question</b>	2
<b>Correct</b>	2
<b>Incorrect</b>	0
<b>Unattempted</b>	0
<b>Marked for review</b>	1
<b>No</b>	4
<b>Domain</b>	Monitor, troubleshoot, and optimize Azure solutions
<b>Total Question</b>	7
<b>Correct</b>	7
<b>Incorrect</b>	0
<b>Unattempted</b>	0
<b>Marked for review</b>	0

<b>No</b>	5
<b>Domain</b>	Develop Azure Platform as a Service Compute Solutions
<b>Total Question</b>	12
<b>Correct</b>	12
<b>Incorrect</b>	0
<b>Unattempted</b>	0
<b>Marked for review</b>	0
<b>No</b>	6
<b>Domain</b>	Implement Azure Security
<b>Total Question</b>	7
<b>Correct</b>	7
<b>Incorrect</b>	0
<b>Unattempted</b>	0
<b>Marked for review</b>	0
<b>Total</b>	Total
<b>All Domain</b>	All Domain
<b>Total Question</b>	55
<b>Correct</b>	55
<b>Incorrect</b>	0
<b>Unattempted</b>	0
<b>Marked for review</b>	1

## Review the Answers

Sorting by

All

Question 1

Correct

Domain :Connect to and Consume Azure Services and Third Party Services

A team has created an Index in the Azure Search service. You have to upload data into the Index. You propose the following steps to carry out from your .Net program

- Create a SearchServiceClient object to connect to the search index.
- Create a DataContainer that contains the documents which must be added.
- Create a DataSource instance and set its Container property to the DataContainer.
- Set the DataSource property of the SearchServiceClient
- Does the list of steps fulfil the requirement?

A. Yes

✓ B. No 

**Explanation:**

Answer – B

The correct list of steps is given below as per an example from the Microsoft documentation

In order to push documents into your index using the .NET SDK, you will need to:

1. Create a `SearchIndexClient` object to connect to your search index.
2. Create an `IndexBatch` containing the documents to be added, modified, or deleted.
3. Call the `Documents.Index` method of your `SearchIndexClient` to send the `IndexBatch` to your search index.

For an example on the steps for importing data, one can go to the below link

<https://docs.microsoft.com/en-us/azure/search/search-import-data-dotnet>

---

Ask our Experts

Rate this Question?  

---

**Question 2**

**Correct**

**Domain :Connect to and Consume Azure Services and Third Party Services**

A team has created an Index in the Azure Search service. You have to upload data into the Index. You propose the following steps to carry out from your .Net program

Create a `SearchIndexClient` object to connect to the search index  
Create an `IndexBatch` that contains the documents which must be added.  
Call the `Documents.Index` method of the `SearchIndexClient` and pass the `IndexBatch`.  
Does the list of steps fulfil the requirement?

- ✓ A. Yes 
- B. No

---

**Explanation:**

Answer – A

Yes, this is the correct list of steps as shown below from the Microsoft documentation

In order to push documents into your index using the .NET SDK, you will need to:

1. Create a `SearchIndexClient` object to connect to your search index.
2. Create an `IndexBatch` containing the documents to be added, modified, or deleted.
3. Call the `Documents.Index` method of your `SearchIndexClient` to send the `IndexBatch` to your search index.

For an example on the steps for importing data, one can go to the below link

<https://docs.microsoft.com/en-us/azure/search/search-import-data-dotnet>

Ask our Experts

Rate this Question?  

Question 3

Correct

#### Domain :Connect to and Consume Azure Services and Third Party Services

A team has created an Index in the Azure Search service. You have to upload data into the Index.

You propose the following steps to carry out from your .Net program

Create a `SearchIndexClient` object to connect to the search index.

Create a `DataContainer` that contains the documents which must be added.

Create a `DataSource` instance and set its `Container` property to the `DataContainer`

Call the `Documents.Search` method of the `SearchIndexClient` and pass the `DataSource`.

Does the list of steps fulfil the requirement?

A. Yes

✓ B. No 

#### Explanation:

Answer – B

The correct list of steps is given below as per an example from the Microsoft documentation

In order to push documents into your index using the .NET SDK, you will need to:

1. Create a `SearchIndexClient` object to connect to your search index.
2. Create an `IndexBatch` containing the documents to be added, modified, or deleted.
3. Call the `Documents.Index` method of your `SearchIndexClient` to send the `IndexBatch` to your search index.

For an example on the steps for importing data, one can go to the below link

<https://docs.microsoft.com/en-us/azure/search/search-import-data-dotnet>

Ask our Experts

Rate this Question?  

Question 4

Correct

Domain :Develop for Azure Storage

A developer has created a table in an Azure storage account called "Customers". The data in the table is partitioned by the column `firstname`. The developer needs to write a query that would return all the customers with the first name of "Dave". Which of the following is the right code segment for the `Where` clause?

- ✓ A. `TableQuery.GenerateFilterCondition("PartitionKey", QueryComparisons.Equal, "Dave")` 
- B. `TableQuery.GenerateFilterCondition("PartitionKey", Equals, "Dave")`
- C. `TableQuery.GenerateFilterCondition("firstname", QueryComparisons.Equal, "Dave")`
- D. `TableQuery.GenerateFilterCondition("firstname", Equal, "Dave")`

#### Explanation:

Answer – A

An example of this is given in the Microsoft documentation as shown below. Here you need to search via the partition key and use the `QueryComparisions` clause.

## Retrieve all entities in a partition

To query a table for all entities in a partition, use a [TableQuery](#) object. The following code example specifies a filter for entities where 'Smith' is the partition key. This example prints the fields of each entity in the query results to the console.

```
C#  
  
string.  
count.Parse(  
nnectionString));  
  
eCloudTableClient();  
  
"people" table.  
eople");  
  
ntities where PartitionKey="Smith".  
ustomerEntity>() Where(TableQuery.GenerateFilterCondition("PartitionKey", QueryComparisons.Equal, "Smith"));  
  
y(query))  
  
PartitionKey, entity.RowKey,  
  
Since this is clearly given in the Microsoft documentation, all other options are incorrect
```



For an example on using table storage from .Net, one can go to the below link

<https://docs.microsoft.com/en-us/azure/cosmos-db/table-storage-how-to-use-dotnet>

Ask our Experts

Rate this Question?  

Question 5

Correct

Domain :Develop for Azure Storage

A company is developing an application for a company. The application is making use of CosmosDB as the backend store. The application store patient details for a hospital.

Below are the key requirements for certain modules

**Requirement 1** - The status for the patient must be the most recent. This should be the case even if multiple users in different locations update the patient's records

**Requirement 2** - The health for a patient is recorded by one module. Here it needs to be ensured that the data must be either the current version or a prior version

**Requirement 3** - When the patient is being discharged, all charges should be processed, and the final bill should be processed

You have to ensure that you minimize the latency and any impact on the availability of the solution

Which of the following consistency level would you choose for Requirement1?

- ✓ A. **Strong** 
- B. **Bounded Staleness**
- C. **Consistent Prefix**
- D. **Eventual**

---

#### Explanation:

Answer – A

Here since the requirement is that the patient's record should be the most consistent, there is a need for consistency in data and no staleness. So, we have to choose Strong consistency for this. The Microsoft documentation mentions the following on the consistency level.

- **Strong:** Strong consistency offers a [linearizability](#) guarantee. The reads are guaranteed to return the most recent committed version of an item. A client never sees an uncommitted or partial write. Users are always guaranteed to read the latest committed write.

The other consistency level options will not give you the required level of consistency

For more information on consistency levels, one can go to the below link

<https://docs.microsoft.com/en-us/azure/cosmos-db/consistency-levels>

---

Ask our Experts

Rate this Question?  

Question 6

Correct

Domain :Develop for Azure Storage

A company is developing an application for a company. The application is making use of CosmosDB as the backend store. The application stores patient details for a hospital.

Below are the key requirements for certain modules

**Requirement 1** - The status for the patient must be the most recent. This should be the case even if multiple users in different locations update the patient's records

**Requirement 2** - The health for a patient is recorded by one module. Here it needs to be ensured that the data must be either the current version or a prior version

**Requirement 3** - When the patient is being discharged, all charges should be processed, and the final bill should be processed

You have to ensure that you minimize the latency and any impact on the availability of the solution

Which of the following consistency level would you choose for Requirement2?

A. Strong

- ✓ B. Bounded Staleness 
- C. Consistent Prefix
- D. Eventual

---

#### Explanation:

Answer – B

Here you can have consistency up to a certain level. From the Microsoft documentation, you can see that this consistency level provides a feature of providing consistency up to a certain number of versions of an item.

- **Bounded staleness:** The reads are guaranteed to honor the consistent-prefix guarantee. The reads might lag behind writes by at most "K" versions (that is "updates") of an item or by "t" time interval. When you choose bounded staleness, the "staleness" can be configured in two ways:
  - The number of versions (K) of the item
  - The time interval (t) by which the reads might lag behind the writes

Since this is the best consistency preference as per the requirement, all the options are invalid.

For more information on consistency levels, one can go to the below link

<https://docs.microsoft.com/en-us/azure/cosmos-db/consistency-levels>

---

Ask our Experts

Rate this Question?  

**Question 7****Correct****Domain :Develop for Azure Storage**

A company is developing an application for a company. The application is making use of CosmosDB as the backend store. The application stores patient details for a hospital.

Below are the key requirements for certain modules

**Requirement 1** - The status for the patient must be the most recent. This should be the case even if multiple users in different locations update the patient's records

**Requirement 2** - The health for a patient is recorded by one module. Here it needs to be ensured that the data must be either the current version or a prior version

**Requirement 3** - When the patient is being discharged, all charges should be processed, and the final bill should be processed

You have to ensure that you minimize the latency and any impact on the availability of the solution

Which of the following consistency level would you choose for Requirement3?

- A. **Strong**
- B. **Bounded Staleness**
- C. **Consistent Prefix**
- ✓ D. **Eventual** 

---

**Explanation:**

Answer – D

Since we here we just need to wait for the final charges, we can just wait for all changes to take effect so here the most effective would be Eventual consistency.

The Microsoft documentation mentions the following on the consistency level.

- **Eventual:** There's no ordering guarantee for reads. In the absence of any further writes, the replicas eventually converge.

Since this is the best consistency preference as per the requirement, all the options are invalid.

For more information on consistency levels, one can go to the below link

<https://docs.microsoft.com/en-us/azure/cosmos-db/consistency-levels>

---

Ask our Experts

Rate this Question?  

**Question 8****Correct****Domain :Develop for Azure Storage**[View Case Study](#)

Which of the following will go into Slot1?

- ✓ A. Region 
- B. Phone
- C. Email
- D. ID

**Explanation:**

Answer – A

Since the question states that we will be using Region to load balance the data, we have to use that as the partition key. The Microsoft documentation mentions the following on the partition key design

### Choosing an appropriate PartitionKey

Your choice of **PartitionKey** should balance the need to enable the use of EGTs (to ensure consistency) against the requirement to distribute your entities across multiple partitions (to ensure a scalable solution).

At one extreme, you could store all your entities in a single partition, but this may limit the scalability of your solution and would prevent the table service from being able to load-balance requests. At the other extreme, you could store one entity per partition, which would be highly scalable and which enables the table service to load-balance requests, but that would prevent you from using entity group transactions.

An ideal **PartitionKey** is one that enables you to use efficient queries and that has sufficient partitions to ensure your solution is scalable. Typically, you will find that your entities will have a suitable property that distributes your entities across sufficient partitions.

For more information on Azure Table storage design, one can go to the below link

<https://docs.microsoft.com/en-us/azure/cosmos-db/table-storage-design-guide>

[Ask our Experts](#)Rate this Question?  

Question 9

Correct

Domain :Develop for Azure Storage

[View Case Study](#)

Which of the following will go into Slot2?

- A. Region
- B. Phone
- C. Email
- D. ID

**Explanation:**

Answer – C

You need to have a value for the RowKey. So, if the Phone field has missing values for some of the entities, then we have to choose Email as the RowKey. The Microsoft documentation mentions the following on the Row Key

**RowKey Property** 

The second part of the primary key is the row key, specified by the `RowKey` property. The row key is a unique identifier for an entity within a given partition. Together the `PartitionKey` and `RowKey` uniquely identify every entity within a table.

The row key is a string value that may be up to 1 KB in size.

You must include the `RowKey` property in every insert, update, and delete operation.

For more information on understanding the table service data model, one can go to the below link

<https://docs.microsoft.com/en-us/rest/api/storageservices/understanding-the-table-service-data-model>

[Ask our Experts](#)Rate this Question?  

Question 10

Correct

Domain :Develop for Azure Storage

[View Case Study](#)

Which of the following will go into Slot3?

- ✓ A. CloudTable 
- B. CloudTableClient
- C. TableEntity
- D. TableEntityAdapter

**Explanation:**

Answer – A

Since we are passing in a table parameter, and this would be a reference to our cloud table, we would need to use the CloudTable data type.

An example snippet of code in the Microsoft documentation is given below

## Create a table

This example shows how to create a table if it does not already exist:

C#

```
// Retrieve the storage account from the connection string.  
CloudStorageAccount storageAccount = CloudStorageAccount.Parse(  
    CloudConfigurationManager.GetSetting("StorageConnectionString"));  
  
// Create the table client.  
CloudTableClient tableClient = storageAccount.CreateCloudTableClient();  
  
// Retrieve a reference to the table.  
CloudTable table = tableClient.GetTableReference("people");  
  
// Create the table if it doesn't exist.  
table.CreateIfNotExists();
```

Since this is clearly given in the Microsoft documentation, all other options are incorrect

For more information on using table storage with .Net, one can go to the below link

<https://docs.microsoft.com/en-us/azure/cosmos-db/table-storage-how-to-use-dotnet>

---

Ask our Experts

Rate this Question?  

Question 11

Correct

Domain :Develop for Azure Storage

[View Case Study](#)

Which of the following will go into Slot4?

- A. TableEntity retrieve=TableEntity.Retrieve(p\_partitionkey,p\_rowkey)
- ✓ B. TableOperation retrieve=TableOperation.Retrieve(p\_partitionkey,p\_rowkey) 
- C. TableResult retrieve=TableQuery.Retrieve(p\_partitionkey,p\_rowkey)
- D. TableResultSegment retrieve=TableResult.Retrieve(p\_partitionkey,p\_rowkey)

---

**Explanation:**

Answer – B

If we need to retrieve an entity based on the partition and row key , we will need to use the TableOperation method.

An example snippet of code in the Microsoft documentation is given below

## Retrieve a single entity

You can write a query to retrieve a single, specific entity. The following code uses [TableOperation](#) to specify the customer 'Ben Smith'. This method returns just one entity rather than a collection, and the returned value in [TableResult.Result](#) is a [CustomerEntity](#) object. Specifying both partition and row keys in a query is the fastest way to retrieve a single entity from the Table service.

C#

 Copy

```
// Retrieve the storage account from the connection string.  
CloudStorageAccount storageAccount = CloudStorageAccount.Parse(  
    CloudConfigurationManager.GetSetting("StorageConnectionString"));  
  
// Create the table client.  
CloudTableClient tableClient = storageAccount.CreateCloudTableClient();  
  
// Create the CloudTable object that represents the "people" table.  
CloudTable table = tableClient.GetTableReference("people");  
  
// Create a retrieve operation that takes a customer entity.  
TableOperation retrieveOperation = TableOperation.Retrieve<CustomerEntity>("Smith", "Ben");  
  
// Execute the retrieve operation.  
TableResult retrievedResult = table.Execute(retrieveOperation);
```

Since this is clearly given in the Microsoft documentation, all other options are incorrect

For more information on using table storage with .Net, one can go to the below link

<https://docs.microsoft.com/en-us/azure/cosmos-db/table-storage-how-to-use-dotnet>

Ask our Experts

Rate this Question?  

Question 12

Correct

Domain :Connect to and Consume Azure Services and Third Party Services

A team is developing a software that will generate large data sets. These large data sets will be processed by nodes in the Azure Batch Service. In the program , you have to create the compute nodes for Azure Batch. Which of the following would you do?

- A. In the program, implement a class for TaskAddParameter

- B. In the program, implement a class for JobAddParameter
- ✓ C. In the program, call the method - BatchClient.PoolOperations.CreatePool 
- D. In the program, call the method - BatchClient.VMOperations.CreateVM

---

**Explanation:**

Answer – C

The Microsoft documentation clearly gives the command for creating a pool of Virtual Machines as shown below

```
private static void CreateBatchPool(BatchClient batchClient, VirtualMachineConfiguration vmConfiguration)
{
    try
    {
        CloudPool pool = batchClient.PoolOperations.CreatePool(
            poolId: PoolId,
            targetDedicatedComputeNodes: PoolNodeCount,
            virtualMachineSize: PoolVMSize,
            virtualMachineConfiguration: vmConfiguration);

        pool.Commit();
    }
    ...
}
```

The Batch client class and the CreatePool method are used to create a pool of instances

Since this is clearly given in the Microsoft documentation, all other options are incorrect

For more information on using .Net to work with Batch clients, one can go to the below link

<https://docs.microsoft.com/en-us/azure/batch/quick-run-dotnet>

---

Ask our Experts

Rate this Question?  

Question 13

Marked as review    Correct

Domain :Develop Azure Infrastructure as a Service Compute Solutions

A team is developing container-based applications that need to be deployed to a Kubernetes cluster in Azure. You have to create the cluster and ensure the services are running as desired. Which of the following commands would you execute? Choose 4 answers from the options given below

- ✓ A. az aks create ✓
- ✓ B. az group create ✓
- ✓ C. kubectl apply ✓
- D. az appservice plan create
- ✓ E. az aks get-credentials ✓

### Explanation:

Answer – A, B, C and E

An example of the steps which need to be followed are given below

The first command is used to create a resource group in Azure. An example is given below

```
PS Azure:> az group create --name whizlabs-rg-new --location eastus
{
  "id": "/subscriptions/baaa99b3-1d19-4c5e-90e1-39d55de5fc6e/resourceGroups/whizlabs-rg-new",
  "location": "eastus",
  "managedBy": null,
  "name": "whizlabs-rg-new",
  "properties": {
    "provisioningState": "Succeeded"
  },
  "tags": null,
  "type": null
}
```

The next command is used to create the Kubernetes cluster. An example is given below

```
PS Azure:> az aks create --resource-group whizlabs-rg-new --name whizlabsclusternew --node-count 2 --enable-addons monitoring --generate-ssh-keys
SSH key files '/home/dcd421e6-6d39-4827-90dc-e641bbe7/.ssh/id_rsa' and '/home/dcd421e6-6d39-4827-90dc-e641bbe7/.ssh/id_rsa.pub' have been generated under ~/.ssh to allow SSH access to the VM. If using machines without permanent storage like Azure Cloud Shell without an attached file share, back up your keys to a safe location
Finished service principal creation[########################################] 100.0000%
```

Next you need to merge the Kubernetes cluster credentials. An example is given below

```
PS Azure:\> az aks get-credentials --resource-group whizlabs-rg-new --name whizlabsclusternew
Merged "whizlabsclusternew" as current context in /home/dcd421e6-6d39-4827-90dc-e641bbe7/.kube/config
Azure:\>
PS Azure:\>
```

And then finally you can use "kubectl apply" to deploy your application onto the cluster

Option D is incorrect since we don't need an App service plan for a Kubernetes cluster.

For more information on using the Azure CLI to work with Kubernetes clusters, one can go to the below link

<https://docs.microsoft.com/en-us/azure/aks/kubernetes-walkthrough>

Ask our Experts

Rate this Question?

Question 14

Correct

Domain :Monitor, troubleshoot, and optimize Azure solutions

A company has a web application that has been deployed using the Azure Web App service. The current service plan being used is D1. It needs to be ensured that the application infrastructure can automatically scale when the CPU load reaches 85 percent. You also have to ensure costs are minimized. Which of the following steps would you implement to achieve the requirements? Choose 4 answers from the options given below

- ✓ A. Enable autoscaling on the Web application
- ✓ B. Configure a scale condition
- ✓ C. Configure the web application to use the Standard App Service Plan
- D. Configure the web application to use the Premium App Service Plan
- ✓ E. Add a scale rule.

#### Explanation:

Answer – A, B, C and E

Since the app service plan being used is D1, that means this is the Shared Service Plan as shown below. And this plan does not have support for Autoscaling



## Dev / Test

For less demanding workloads



## Production

For most production workloads



## Isolated

Advanced networking and scale

## Recommended pricing tiers

F1	Shared infrastructure 1 GB memory 60 minutes/day compute Loading...	D1	Shared infrastructure 1 GB memory 240 minutes/day compute Loading...	B1	100 total ACU 1.75 GB memory A-Series compute equivalent Loading...
----	------------------------------------------------------------------------------	----	-------------------------------------------------------------------------------	----	------------------------------------------------------------------------------

[See additional options](#)

## Included features

Every app hosted on this App Service plan will have access to these features:



## Custom domains

Configure and purchase custom domain names.

## Included hardware

Every instance of your App Service plan will include the following hardware configuration:



## Azure Compute Units (ACU)

Dedicated compute resources used to run applications deployed in the App Service Plan. [Learn more](#)

## Memory

Memory available to run applications deployed and running in the App Service plan.



## Storage

1 GB disk storage shared by all apps deployed in the App Service plan.

Step 1) We have to scale up to at least the Standard App service plan.

Once this is done, you can now see the ability to enable Autoscale when you go to the Scale out section for the Azure Web App

whizlabsapp - Scale out (App Service plan)  
App Service

Search (Ctrl+ /)

Save Discard Disable autoscale Refresh

Configure Run history JSON Notify

Override condition

Instance count 1

Your autoscale configuration is disabled. To reinstate your configuration, enable autoscale.

Enable autoscale

Step 2) Next you add a scale condition and a rule for autoscaling based on a CPU threshold.

out (App Service plan)

### whizlabsapp - Scale out (App Service plan)

App Service

 Search (Ctrl+/  
Save Discard Disable autoscale Refresh
[Configure](#) [Run history](#) [JSON](#) [Notify](#)

- [Deployment credentials](#)
- [Deployment slots](#)
- [Deployment Center](#)

**Settings**

- [Application settings](#)
- [Configuration \(Preview\)](#)
- [Authentication / Authorization](#)
- [Application Insights](#)
- [Identity](#)
- [Backups](#)
- [Custom domains](#)
- [SSL settings](#)
- [Networking](#)
- [Scale up \(App Service plan\)](#)
- [Scale out \(App Service plan\)](#)
- [WebJobs](#)
- [Push](#)
- [MySQL In App](#)
- [Properties](#)
- [Locks](#)

**Configure** Run history JSON Notify

\* Autoscale setting name: CPU 1

Resource group: whizlabs-rg

**Default** Auto created scale condition Edit

Delete warning: The very last or default recurrence rule cannot be deleted. Instead, you can disable it.

Scale mode:  Scale based on a metric  Scale to a specific instance count 2  
Scale out and scale in your instances based on metric. For example: 'Add a rule that inc above 70%'

Rules: + Add a rule

Instance limits: Minimum 1, Maximum 2, Default 1

Schedule: This scale condition is executed when none of the other scale condition(s) match

+ Add a scale condition

**Scale rule**

Metric source: Current resource (demoplan) ▼

Resource type: App Service plans ▼

Resource: demoplan ▼

**Criteria**

\* Time aggregation: Average ▼

\* Metric name: CPU Percentage ▼  
1 minute time grain

\* Time grain statistic: Average ▼

\* Operator: Greater than ▼

\* Threshold: 85 ✓

\* Duration (in minutes): 10

Action:

\* Operation: Add

Option D is incorrect since the Premium app service plan would be a more expensive option.

For more information on Azure Web App Autoscaling, one can go to the below link

<https://blogs.msdn.microsoft.com/benjaminperkins/2017/07/26/how-to-configure-auto-scaling-for-an-azure-app-service-with-powershell/>

Ask our Experts

Rate this Question? 😊 😢

**Question 15**

**Correct**

**Domain :Monitor,troubleshoot, and optimize Azure solutions**

A company is planning on using the Azure CDN service to distribute static images.

Below is a set of steps in a random order which would be followed by the CDN service.

1. The Origin server will return an image to the Edge server in the Point of Presence. The edge server will cache the image and return the image to the user
  2. If no edge server in the Point of Presence has an image in the cache, the Point of Presence will request the image from the origin server.
  3. A user requests an image from the CDN URL. The DNS routes the request to the best performing Point of Presence location
  4. Subsequent requests for the image may be directed to the same Point of Presence. The Point of Presence will return the image if the TTL has not expired.
- Which of the following is the correct process of how the Content Delivery service would distribute the images?

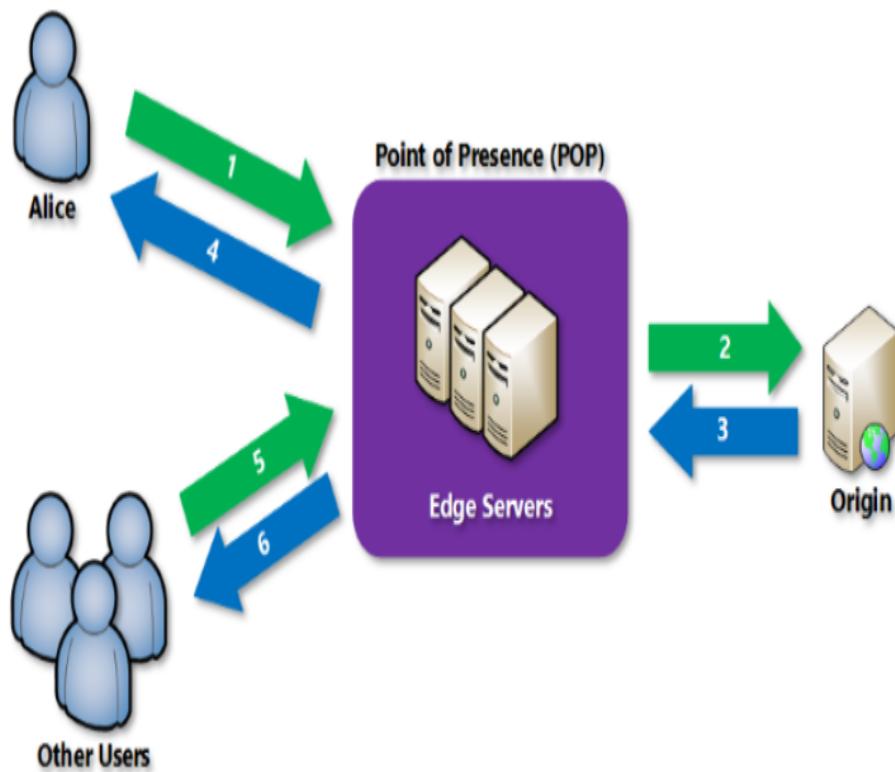
- ✓ A. **3,2,1,4** 
- B. **2,1,3,4**
- C. **1,2,3,4**
- D. **4,3,2,1**

---

#### Explanation:

Answer - A

An example of this is given in the Microsoft documentation



1. A user (Alice) requests a file (also called an asset) by using a URL with a special domain name, such as <endpoint name>.azureedge.net. This name can be an endpoint hostname or a custom domain. The DNS routes the request to the best performing POP location, which is usually the POP that is geographically closest to the user.
2. If no edge servers in the POP have the file in their cache, the POP requests the file from the origin server. The origin server can be an Azure Web App, Azure Cloud Service, Azure Storage account, or any publicly accessible web server.
3. The origin server returns the file to an edge server in the POP.
4. An edge server in the POP caches the file and returns the file to the original requestor (Alice). The file remains cached on the edge server in the POP until the time-to-live (TTL) specified by its HTTP headers expires. If the origin server didn't specify a TTL, the default TTL is seven days.
5. Additional users can then request the same file by using the same URL that Alice used, and can also be directed to the same POP.
6. If the TTL for the file hasn't expired, the POP edge server returns the file directly from the cache. This process results in a faster, more responsive user experience.

Since the steps are clearly mentioned, all other options are incorrect

For more information on Azure Web App Autoscaling, one can go to the below link

<https://docs.microsoft.com/en-us/azure/cdn/cdn-overview>

[Ask our Experts](#)Rate this Question?  

Question 16

Correct

**Domain :Connect to and Consume Azure Services and Third Party Services**

A team needs to manage several Logic Apps. There is a need to change definitions, add new logic and optimize the applications on a regular basis. You need to ensure that you use the right tools for the right purpose.

Which of the following would you use to edit B2B workflows?

- A. Logic Apps Designer
- B. Code View Editor
- C. Enterprise Integration Pack 
- D. API Connections

**Explanation:**

Answer – C

This is given in the Microsoft documentation

## Receive B2B data with Azure Logic Apps and Enterprise Integration Pack

07/08/2016 • 2 minutes to read • Contributors 

After you create an integration account that has partners and agreements, you are ready to create a business to business (B2B) workflow for your logic app with the [Enterprise Integration Pack](#).

Since this is clearly given, all other options are incorrect

For more information on enterprise integration with B2B, one can go to the below link

<https://docs.microsoft.com/en-us/azure/logic-apps/logic-apps-enterprise-integration-b2b>

[Ask our Experts](#)

Rate this Question?

Question 17

Correct

**Domain :Connect to and Consume Azure Services and Third Party Services**

A team needs to manage several Logic Apps. There is a need to change definitions, add new logic and optimize the applications on a regular basis. You need to ensure that you use the right tools for the right purpose.

Which of the following would you use to edit definitions in JSON?

- A. Logic Apps Designer
- B. Code View Editor
- C. Enterprise Integration Pack
- D. API Connections

**Explanation:**

Answer – B

You can use the App Code view of Azure Logic Apps as shown below.

```

1   "definition": {
2     "$schema": "https://schema.management.azure.com/providers/Microsoft.Logic/schemas/2016-06-01/workflowdefinition.json#",
3     "actions": {},
4     "contentVersion": "1.0.0.0",
5     "outputs": {},
6     "parameters": {},
7     "triggers": {}
8   }
9
10 }

```

Since this is evident from the implementation, all other options are invalid.

For more information on authoring definitions in Azure Logic Apps, one can go to the below link

<https://docs.microsoft.com/en-us/azure/logic-apps/logic-apps-author-definitions>

---

Ask our Experts

Rate this Question?  

Question 18

Correct

**Domain :Connect to and Consume Azure Services and Third Party Services**

A team needs to manage several Logic Apps. There is a need to change definitions, add new logic and optimize the applications on a regular basis. You need to ensure that you use the right tools for the right purpose.

Which of the following would you use to visually add functionality to the Logic App?

- ✓ A. Logic Apps Designer 
- B. Code View Editor
- C. Enterprise Integration Pack
- D. API Connections

---

**Explanation:**

Answer – A

You can use the Logic Apps Designer as shown below. There you can create and add functionality to the Logic App using the Designer.

The screenshot shows the Azure Logic Apps Designer interface. On the left, a sidebar menu is open under 'Development Tools' with 'Logic app designer' selected (marked with a red circle containing the number 1). The main area displays a video thumbnail titled 'Introducing Azure Logic Apps' with a play button. Below the video, four icons represent common triggers: 'When a message is received in a Service Bus queue', 'When a HTTP request is received', 'When an Event Grid resource event occurs', and 'Recurrence'. To the right of these triggers, a section titled 'Building integration solutions is easier than ever' describes Logic Apps and lists three benefits: creating business processes visually, integrating with SaaS and enterprise applications, and unlocking value from on-premises and cloud applications.

Since this is evident from the implementation, all other options are invalid.

For more information on an example on using Azure Logic Apps, one can go to the below link

<https://docs.microsoft.com/en-us/azure/logic-apps/quickstart-create-first-logic-app-workflow>

#### Ask our Experts

Rate this Question?

Question 19

Correct

#### Domain :Connect to and Consume Azure Services and Third Party Services

A team is developing a project management service by using ASP.NET. One of the modules of the service needs to allow users to search for keywords in the index data hosted in Azure Search. You need to implement code that creates the object which is used to create indexes in the Azure Search service. Which of the following objects would you use for this purpose? Choose 2 answers from the options given below

- A. **SearchService**

**B. SearchIndexClient**

- ✓ C. **SearchServiceClient**
- ✓ D. **SearchCredentials**

**Explanation:**

Answer – C and D

You have to use the SearchServiceClient and use the Indexes property as shown below in the Microsoft documentation. And the other object you would use is the SearchCredentials object which has the credentials to connect to the Search service.

```
C#  Copy

private static SearchServiceClient CreateSearchServiceClient(IConfigurationRoot configuration)
{
    string searchServiceName = configuration["SearchServiceName"];
    string adminApiKey = configuration["SearchServiceAdminApiKey"];

    SearchServiceClient serviceClient = new SearchServiceClient(searchServiceName, new SearchCredentials(adminApiKey));
    return serviceClient;
}
```

SearchServiceClient has an Indexes property. This property provides all the methods you need to create, list, update, or delete Azure Search indexes.

Since this is clearly given in the Microsoft documentation, all other options are incorrect

For more information on creating an index in .Net, one can go to the below link

<https://docs.microsoft.com/en-us/azure/search/search-create-index-dotnet>

Ask our Experts

Rate this Question?

Question 20

Correct

Domain :Develop Azure Platform as a Service Compute Solutions

A company needs to develop a script that will do the following

Create an Azure Web App

Create the Web App service plan

Ensure automatic deployment of code from Github

The following variables are in place

Variable Name	Variable value
\$gitrepo	<a href="https://github.com/whizlabs/app">https://github.com/whizlabs/app</a>
\$webappname	whizlabsapp

You have to complete the following script

```
az group create --location westeurope --name "whizlabs-rg"
```

Slot1

```
--name $webappname --resource-group "whizlabs-rg" --sku FREE
```

Slot2

```
--name $webappname --resource-group "whizlabs-rg" --plan $webappname
```

Slot3

```
source config --name $webappname --resource-group "whizlabs-rg" \
```

Slot4

```
$gitrepo --branch master --manual-integration
```

Which of the following would go into Slot1?

- A. az webapp create
- B. az appservice plan create
- C. az webapp deployment
- D. az group assign

#### Explanation:

Answer – B

This is given as an example in the Microsoft documentation

Azure CLI

Copy

Try It

```
#!/bin/bash

# Replace the following URL with a public GitHub repo URL
gitrepo=https://github.com/Azure-Samples/php-docs-hello-world
webappname=mywebapp$RANDOM

# Create a resource group.
az group create --location westeurope --name myResourceGroup

# Create an App Service plan in `FREE` tier.
az appservice plan create --name $webappname --resource-group myResourceGroup --sku FREE

# Create a web app.
az webapp create --name $webappname --resource-group myResourceGroup --plan $webappname

# Deploy code from a public GitHub repository.
az webapp deployment source config --name $webappname --resource-group myResourceGroup \
--repo-url $gitrepo --branch master --manual-integration

# Copy the result of the following command into a browser to see the web app.
echo http://$webappname.azurewebsites.net
```

Ask our Experts

Rate this Question?

Question 21

Correct

#### Domain :Develop Azure Platform as a Service Compute Solutions

A company needs to develop a script that will do the following

Create an Azure Web App

Create the Web App service plan

Ensure automatic deployment of code from Github

The following variables are in place

Variable Name

Variable value

\$gitrepo

<https://github.com/whizlabs/app>

\$webappname

whizlabsapp

You have to complete the following script

az group create --location westeurope --name "whizlabs-rg"

Slot1

--name \$webappname --resource-group "whizlabs-rg" --sku FREE

Slot2

--name \$webappname --resource-group "whizlabs-rg" --plan \$webappname

Slot3

source config --name \$webappname --resource-group "whizlabs-rg" \

Slot4

\$gitrepo --branch master --manual-integration

Which of the following would go into Slot2?

- ✓ A. az webapp create 
- B. az appservice plan create
- C. az webapp deployment
- D. az group assign

#### Explanation:

Answer – A

This is given as an example in the Microsoft documentation

```
#!/bin/bash

# Replace the following URL with a public GitHub repo URL
gitrepo=https://github.com/Azure-Samples/php-docs-hello-world
webappname=mywebapp$RANDOM

# Create a resource group.
az group create --location westeurope --name myResourceGroup

# Create an App Service plan in `FREE` tier.
az appservice plan create --name $webappname --resource-group myResourceGroup --sku FREE

# Create a web app.
az webapp create --name $webappname --resource-group myResourceGroup --plan $webappname

# Deploy code from a public GitHub repository.
az webapp deployment source config --name $webappname --resource-group myResourceGroup \
--repo-url $gitrepo --branch master --manual-integration

# Copy the result of the following command into a browser to see the web app.
echo http://$webappname.azurewebsites.net
```

Since this is clearly given in the documentation, all other options are incorrect

For more information on using the CLI for deployment, one can go to the below link

<https://docs.microsoft.com/en-us/azure/app-service/scripts/cli-deploy-github>

Ask our Experts

Rate this Question?

Question 22

Correct

#### Domain :Develop Azure Platform as a Service Compute Solutions

A company needs to develop a script that will do the following

Create an Azure Web App

Create the Web App service plan

Ensure automatic deployment of code from Github

The following variables are in place

Variable Name

Variable value

\$gitrepo

<https://github.com/whizlabs/app>

\$webappname

whizlabsapp

You have to complete the following script

az group create --location westeurope --name "whizlabs-rg"

Slot1

--name \$webappname --resource-group "whizlabs-rg" --sku FREE

Slot2

--name \$webappname --resource-group "whizlabs-rg" --plan \$webappname

Slot3

source config --name \$webappname --resource-group "whizlabs-rg" \

Slot4

\$gitrepo --branch master --manual-integration

Which of the following would go into Slot3?

- A. az webapp create
- B. az appservice plan create
- C. az webapp deployment 
- D. az group assign

#### Explanation:

Answer – C

This is given as an example in the Microsoft documentation

Azure CLI

 Copy  Try It

```

#!/bin/bash

# Replace the following URL with a public GitHub repo URL
gitrepo=https://github.com/Azure-Samples/php-docs-hello-world
webappname=mywebapp$RANDOM

# Create a resource group.
az group create --location westeurope --name myResourceGroup

# Create an App Service plan in `FREE` tier.
az appservice plan create --name $webappname --resource-group myResourceGroup --sku FREE

# Create a web app.
az webapp create --name $webappname --resource-group myResourceGroup --plan $webappname

# Deploy code from a public GitHub repository.
az webapp deployment source config --name $webappname --resource-group myResourceGroup \
--repo-url $gitrepo --branch master --manual-integration

# Copy the result of the following command into a browser to see the web app.
echo http://$webappname.azurewebsites.net

```

Since this is clearly given in the documentation, all other options are incorrect

For more information on using the CLI for deployment, one can go to the below link

<https://docs.microsoft.com/en-us/azure/app-service/scripts/cli-deploy-github>

Ask our Experts

Rate this Question?  

Question 23

Correct

Domain :Develop Azure Platform as a Service Compute Solutions

A company needs to develop a script that will do the following

Create an Azure Web App

Create the Web App service plan

Ensure automatic deployment of code from Github

The following variables are in place

Variable Name
---------------

Variable value
----------------

\$gitrepo
-----------

https://github.com/whizlabs/app
---------------------------------

\$webappname
--------------

whizlabsapp
-------------

You have to complete the following script

```
az group create --location westeurope --name "whizlabs-rg"
```

Slot1

```
--name $webappname --resource-group "whizlabs-rg" --sku FREE
```

Slot2

```
--name $webappname --resource-group "whizlabs-rg" --plan $webappname
```

Slot3

```
source config --name $webappname --resource-group "whizlabs-rg" \
```

Slot4

```
$gitrepo --branch master --manual-integration
```

Which of the following would go into Slot4?

- ✓ A. --repo-url
- B. --github-deploy
- C. --github-repo
- D. --repo-deploy

---

**Explanation:**

Answer – A

This is given as an example in the Microsoft documentation

Azure CLI

Copy Try It

```
#!/bin/bash

# Replace the following URL with a public GitHub repo URL
gitrepo=https://github.com/Azure-Samples/php-docs-hello-world
webappname=mywebapp$RANDOM

# Create a resource group.
az group create --location westeurope --name myResourceGroup

# Create an App Service plan in `FREE` tier.
az appservice plan create --name $webappname --resource-group myResourceGroup --sku FREE

# Create a web app.
az webapp create --name $webappname --resource-group myResourceGroup --plan $webappname

# Deploy code from a public GitHub repository.
az webapp deployment source config --name $webappname --resource-group myResourceGroup \
--repo-url $gitrepo --branch master --manual-integration

# Copy the result of the following command into a browser to see the web app.
echo http://$webappname.azurewebsites.net
```

Since this is clearly given in the documentation, all other options are incorrect

For more information on using the CLI for deployment, one can go to the below link

<https://docs.microsoft.com/en-us/azure/app-service/scripts/cli-deploy-github>

Ask our Experts

Rate this Question?

Question 24

Correct

Domain :Develop Azure Platform as a Service Compute Solutions

Your team is developing Azure Web jobs. You have to decide on the type of Webjobs to be used for different scenarios.

Runs on all instances that the web app runs on.

Have the ability to restrict the web job to run on a single instance

Supports remote debugging

Ability to run based on a schedule

Which of the following web job type would you choose for the below requirement?

"Runs on all instances that the web app runs on."

A. Triggered

✓ B. Continuous 

C. Scheduled

D. Instance

#### Explanation:

Answer – B

This is given in the Microsoft documentation as shown below

## WebJob types

The following table describes the differences between *continuous* and *triggered* WebJobs.

Continuous	Triggered
Starts immediately when the WebJob is created. To keep the job from ending, the program or script typically does its work inside an endless loop. If the job does end, you can restart it.	Starts only when triggered manually or on a schedule.
Runs on all instances that the web app runs on. You can optionally restrict the WebJob to a single instance.	Runs on a single instance that Azure selects for load balancing.
Supports remote debugging.	Doesn't support remote debugging.

Since this is clearly given in the documentation, all other options are incorrect

For more information on creating web jobs, one can go to the below link

<https://docs.microsoft.com/en-us/azure/app-service/webjobs-create>

---

Ask our Experts

Rate this Question?  

Question 25

Correct

Domain :Develop Azure Platform as a Service Compute Solutions

Your team is developing Azure Web jobs. You have to decide on the type of Webjobs to be used for different scenarios.

Runs on all instances that the web app runs on.

.Have the ability to restrict the web job to run on a single instance

Supports remote debugging

.Ability to run based on a schedule

Which of the following web job type would you choose for the below requirement?

**"Have the ability to restrict the web job to run on a single instance"**

- A. Triggered
- B. Continuous 
- C. Scheduled
- D. Instance

---

**Explanation:**

Answer – B

This is given in the Microsoft documentation as shown below

# WebJob types

The following table describes the differences between *continuous* and *triggered* WebJobs.

Continuous	Triggered
Starts immediately when the WebJob is created. To keep the job from ending, the program or script typically does its work inside an endless loop. If the job does end, you can restart it.	Starts only when triggered manually or on a schedule.
Runs on all instances that the web app runs on. You can optionally restrict the WebJob to a single instance.	Runs on a single instance that Azure selects for load balancing.
Supports remote debugging.	Doesn't support remote debugging.

Since this is clearly given in the documentation, all other options are incorrect

For more information on creating web jobs, one can go to the below link

<https://docs.microsoft.com/en-us/azure/app-service/webjobs-create>

Ask our Experts

Rate this Question?  

Question 26

Correct

Domain :Develop Azure Platform as a Service Compute Solutions

Your team is developing Azure Web jobs. You have to decide on the type of Webjobs to be used for different scenarios.

Runs on all instances that the web app runs on.

.Have the ability to restrict the web job to run on a single instance

Supports remote debugging

.Ability to run based on a schedule

Which of the following web job type would you choose for the below requirement?

**"Supports remote debugging"**

- A. Triggered
- ✓ B. Continuous 
- C. Scheduled
- D. Instance

**Explanation:**

Answer – B

This is given in the Microsoft documentation as shown below

## WebJob types

The following table describes the differences between *continuous* and *triggered* WebJobs.

Continuous	Triggered
Starts immediately when the WebJob is created. To keep the job from ending, the program or script typically does its work inside an endless loop. If the job does end, you can restart it.	Starts only when triggered manually or on a schedule.
Runs on all instances that the web app runs on. You can optionally restrict the WebJob to a single instance.	Runs on a single instance that Azure selects for load balancing.
Supports remote debugging.	Doesn't support remote debugging.

Since this is clearly given in the documentation, all other options are incorrect

For more information on creating web jobs, one can go to the below link

<https://docs.microsoft.com/en-us/azure/app-service/webjobs-create>

Ask our Experts

Rate this Question?  

**Question 27****Correct****Domain :Develop Azure Platform as a Service Compute Solutions**

Your team is developing Azure Web jobs. You have to decide on the type of Webjobs to be used for different scenarios.

Runs on all instances that the web app runs on.

Have the ability to restrict the web job to run on a single instance

Supports remote debugging

Ability to run based on a schedule

Which of the following web job type would you choose for the below requirement?

**"Ability to run based on a schedule"**

- ✓ A. Triggered 
- B. Continuous
- C. Scheduled
- D. Instance

---

**Explanation:**

Answer – A

This is given in the Microsoft documentation as shown below

# WebJob types

The following table describes the differences between *continuous* and *triggered* WebJobs.

Continuous	Triggered
Starts immediately when the WebJob is created. To keep the job from ending, the program or script typically does its work inside an endless loop. If the job does end, you can restart it.	Starts only when triggered manually or on a schedule.
Runs on all instances that the web app runs on. You can optionally restrict the WebJob to a single instance.	Runs on a single instance that Azure selects for load balancing.
Supports remote debugging.	Doesn't support remote debugging.

Since this is clearly given in the documentation, all other options are incorrect

For more information on creating web jobs, one can go to the below link

<https://docs.microsoft.com/en-us/azure/app-service/webjobs-create>

Ask our Experts

Rate this Question?  

Question 28

Correct

Domain :Monitor, troubleshoot, and optimize Azure solutions

A company is developing a web site. They are planning on deploying the web site to Azure. There is a requirement to ensure that the web site remains available when it experiences high volumes of traffic. You need to minimize on cost. Which of the following would you consider from a deployment perspective?

- A. Deploy the website to an App Service that uses the Shared service tier. Configure the App Service plan to automatically scale when the CPU load is high.

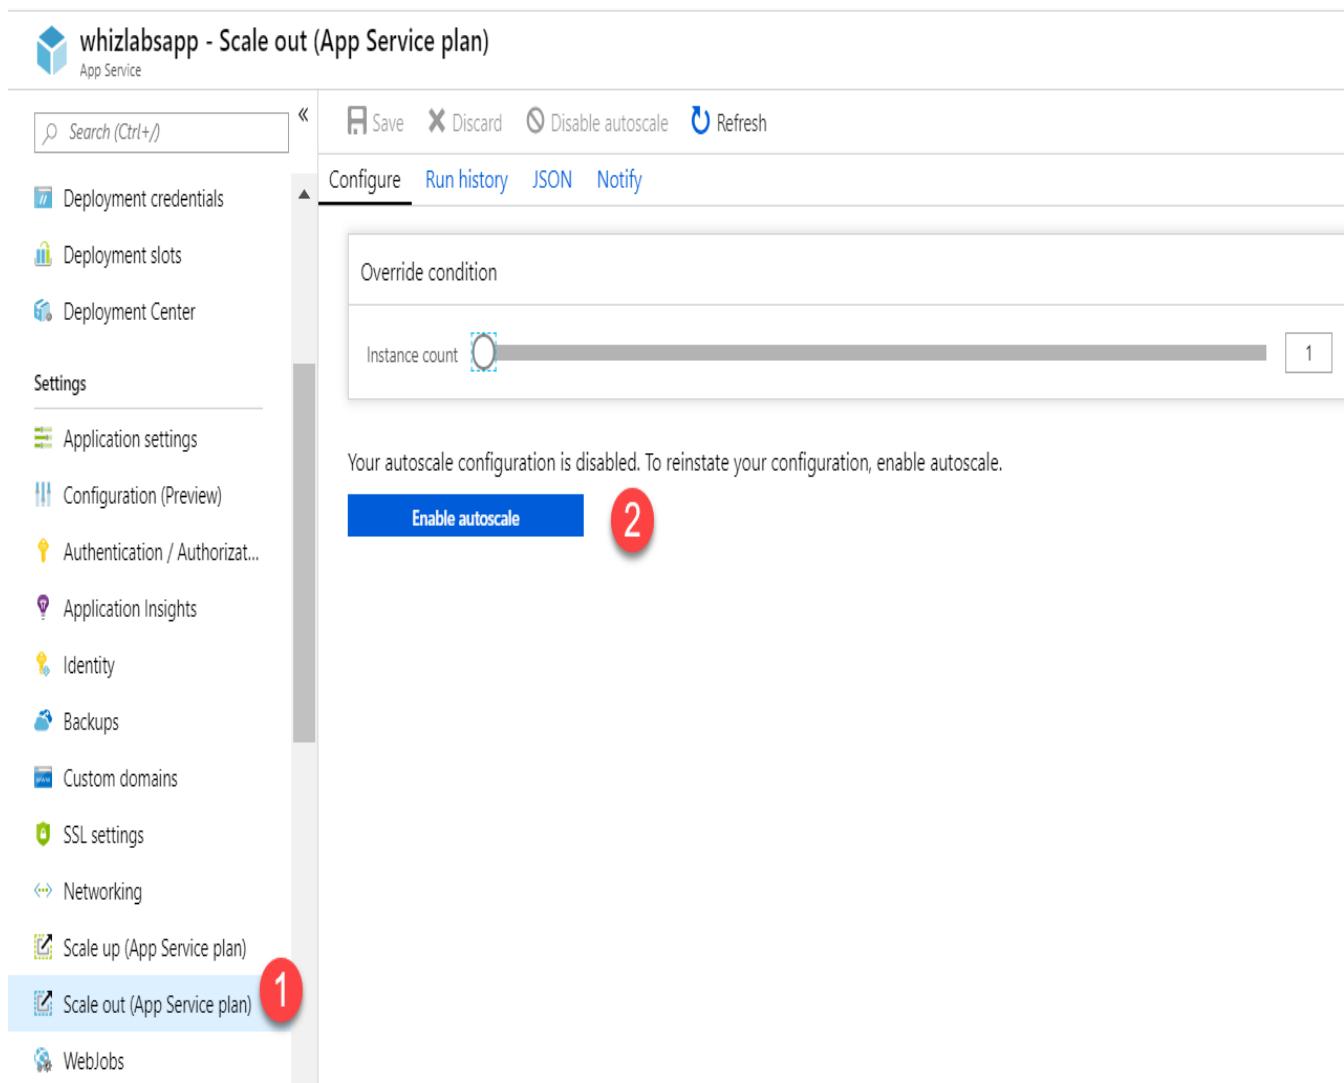
- B. Deploy the website to a virtual machine. Configure the virtual machine to automatically scale when the CPU load is high.
  
- ✓ C. Deploy the website to an App Service that uses the Standard service tier. ✓
  
- D. Configure the App Service plan to automatically scale when the CPU load is high.
  
- D. Deploy the website to a virtual machine. Configure a Scale Set to increase the virtual machine instance count when the CPU load

**Explanation:**

Answer – C

Web Apps deployed to the Standard App Service Plan have the ability to scale up based on demand.

Below is the sample implementation snapshots.



The screenshot shows the Azure portal interface for managing an App Service plan named 'whizlabsapp'. The 'Configure' tab is active. In the main pane, there is a section titled 'Override condition' with a slider for 'Instance count' set to 1. Below this, a message states: 'Your autoscale configuration is disabled. To reinstate your configuration, enable autoscale.' A blue button labeled 'Enable autoscale' is visible. On the left sidebar, under the 'Settings' category, the 'Scale out (App Service plan)' link is highlighted with a red circle containing the number '1'. A red circle containing the number '2' is placed over the 'Enable autoscale' button.

You can add a scale condition and a rule for autoscaling based on a CPU threshold.

out (App Service plan)

### whizlabsapp - Scale out (App Service plan)

App Service

Save Discard Disable autoscale Refresh

Configure Run history JSON Notify

\* Autoscale setting name: CPU (1)

Resource group: whizlabs-rg

**Default** Auto created scale condition

Delete warning: The very last or default recurrence rule cannot be deleted. Instead, you can disable it.

Scale mode: Scale based on a metric (2) Scale to a specific instance count

Scale out and scale in your instances based on metric. For example: 'Add a rule that inc above 70%'

Rules: It is recommended to have at least one scale in rule

+ Add a rule

Instance limits: Minimum 1, Maximum 2, Default 1

Schedule: This scale condition is executed when none of the other scale condition(s) match

+ Add a scale condition

Scale rule

Metric source: Current resource (demoplan)

Resource type: App Service plans

Resource: demoplan

Criteria

\* Time aggregation: Average

\* Metric name: CPU Percentage

1 minute time grain

\* Time grain statistic: Average

\* Operator: Greater than

\* Threshold: 85 ✓

\* Duration (in minutes): 10

Action

\* Operation

Add

Option A is incorrect since the Shared App Service plan does not have the option to scale the application based on demand.

Option B is incorrect since you need to scale amongst multiple instances

Option D is incorrect since this would not be cost effective

For more information on Azure Web App Autoscaling, one can go to the below link

<https://blogs.msdn.microsoft.com/benjaminperkins/2017/07/26/how-to-configure-auto-scaling-for-an-azure-app-service-with-powershell/>

Ask our Experts

Rate this Question? 😊 😐

Question 29

Correct

**Domain :Develop Azure Infrastructure as a Service Compute Solutions**

A developer is creating code to implement the Azure batch service. You have to implement the method to submit a job to the Batch service. Which method would you use?

Choose 2 Options.

- A. `JobOperations.CreateJob()` 
- B. `CloudJob.Enable(IEnumerable)`
- C. `CloudJob.CommitAsync(IEnumerable, CancellationToken)` 
- D. `JobOperations.EnableJob(String, Enumerable)`
- E. `JobOperations.EnableJobAsync(String, IEnumerable, CancellationToken)`

---

**Explanation:**

Answer – A and C

commit :

URL <https://docs.microsoft.com/en-us/dotnet/api/microsoft.azure.batch.cloudjob.commit?view=azure-dotnet>

Commits this CloudJob to the Azure Batch service. This is a blocking operation.

commitAsync :

URL <https://docs.microsoft.com/en-us/dotnet/api/microsoft.azure.batch.cloudjob.commitasync?view=azure-dotnet>

Commits this CloudJob to the Azure Batch service. The commit operation runs asynchronously.

This is given as an example in the Microsoft documentation

## Create a Batch job

A Batch job is a logical grouping of one or more tasks. A job includes settings common to the tasks, such as priority and the pool to run tasks on. The app uses the [BatchClient.JobOperations.CreateJob](#) method to create a job on your pool.

The [Commit](#) method submits the job to the Batch service. Initially the job has no tasks.

C#

 Copy

```
try
{
    CloudJob job = batchClient.JobOperations.CreateJob();
    job.Id = JobId;
    job.PoolInformation = new PoolInformation { PoolId = PoolId };

    job.Commit();
}
...
...
```

Since this is clearly given in the Microsoft documentation, all other options are incorrect

For more information on an example of using Azure Batch with .Net, one can go to the below link

<https://docs.microsoft.com/en-us/azure/batch/quick-run-dotnet>

Ask our Experts

Rate this Question?  

Question 30

Correct

Domain :Implement Azure Security

A development is deploying a web application to Azure. The Web application uses Azure Active Directory for authentication. There is a requirement to implement multifactor authentication for the application. Which of the following needs to be done to fulfil this requirement? Choose 2 answers from the options given below

- ✓ A. In Azure AD, create a new conditional access policy. 
- B. In Azure AD, enable application proxy
- C. Configure the website to use Azure AD B2C.
- D. In Azure AD conditional access, enable the baseline policy.
- ✓ E. Upgrade to Azure AD Premium. 

---

#### Explanation:

Answer - A and E

The Microsoft documentation mentions the following

Conditional access policies can be used to ensure that Multi-factor is implemented for users.

The Microsoft documentation also has a quick start tutorial on how to implement MFA for apps using conditional access policies.

## Quickstart: Require MFA for specific apps with Azure Active Directory conditional access

01/30/2019 • 4 minutes to read • Contributors  all

To simplify the sign-in experience of your users, you might want to allow them to sign in to your cloud apps using a user name and a password. However, many environments have at least a few apps for which it is advisable to require a stronger form of account verification, such as multi-factor authentication (MFA). This might be, for example true, for access to your organization's email system or your HR apps. In Azure Active Directory (Azure AD), you can accomplish this goal with a conditional access policy.

This quickstart shows how to configure an [Azure AD conditional access policy](#) that requires multi-factor authentication for a selected cloud app in your environment.

To use conditional access policies, you need to have Azure AD Premium licensing.

## License requirements for using conditional access

Using conditional access requires an Azure AD Premium license. To find the right license for your requirements, see [Comparing generally available features of the Free, Basic, and Premium editions](#).

Option B is incorrect since you need to use conditional access policies and not application proxy

Option C is incorrect since this does not have the feature for Multi-factor authentication

Option D is incorrect since the baseline is already in place and only allows MFA for administrators

For more information on conditional access policies, one can go to the below link

<https://docs.microsoft.com/en-us/azure/active-directory/conditional-access/overview>

---

Ask our Experts

Rate this Question?

Question 31

Correct

Domain :Implement Azure Security

A company is setting role definitions for various departments in their IT department. Below are some of the requirements

The development teams should be able to communicate with Microsoft support

The IT customer center team should be able to view Azure resources and create support tickets

Custom role definitions need to be in place for the various requirements based on existing role definitions.

Which of the following powershell command would you use to create the custom role?

- ✓ A. `Get-AzRoleDefinition -Name "Reader" | ConvertTo-Json Out-File C:\Sample.json`
- B. `Set-AzRoleDefinition -Name "Reader" | ConvertTo-Json Out-File C:\Sample.json`
- C. `Update-AzRoleDefinition -Name "Reader" | ConvertTo-Json Out-File C:\Sample.json`
- D. `Update-AzRoleDefinition -Name "Reader"`

---

#### Explanation:

Answer – A

First you have to get an existing role definition using the Get-AzRoleDefinition command. An example from the Microsoft documentation is given below

## Create a custom role with the PSRoleDefinition object

When you use PowerShell to create a custom role, you can use one of the [built-in roles](#) as a starting point or you can start from scratch. The first example in this section starts with a built-in role and then customizes it with more permissions. Edit the attributes to add the `Actions`, `NotActions`, or `AssignableScopes` that you want, and then save the changes as a new role.

The following example starts with the [Virtual Machine Contributor](#) built-in role to create a custom role named *Virtual Machine Operator*. The new role grants access to all read operations of *Microsoft.Compute*, *Microsoft.Storage*, and *Microsoft.Network* resource providers and grants access to start, restart, and monitor virtual machines. The custom role can be used in two subscriptions.

```
Azure PowerShell Copy
$role = Get-AzRoleDefinition "Virtual Machine Contributor"
```

Since this is clearly shown in the documentation, all other options are incorrect

For more information on creating custom roles, one can go to the below link

<https://docs.microsoft.com/en-us/azure/role-based-access-control/custom-roles-powershell>

[Ask our Experts](#)

Rate this Question?  

Question 32

Correct

Domain :Implement Azure Security

A company is setting role definitions for various departments in their IT department. Below are some of the requirements

The development teams should be able to communicate with Microsoft support

The IT customer center team should be able to view Azure resources and create support tickets

Custom role definitions need to be in place for the various requirements based on existing role definitions.

Which of the following action would be applicable to allow a team to contact Microsoft support?

- A. `"/read.Microsoft.support"`

B.    \*\*\*

- ✓ C. "Microsoft.Support/\*" 
- D. /\*/Microsoft.Support/\*"

---

**Explanation:**

Answer – C

If you look at the built-in role and see an example for actions to allow support to raise Microsoft support tickets, you can see the required action

---

Microsoft.Support/\*

Create and manage support tickets

---

Since this is clearly shown in the documentation, all other options are incorrect

For more information on the built-in roles, one can go to the below link

<https://docs.microsoft.com/en-us/azure/role-based-access-control/built-in-roles>

---

Ask our Experts

Rate this Question?  

Question 33

Correct

Domain :Monitor, troubleshoot, and optimize Azure solutions

A team is developing an ASP.NET Core application. The Application needs to log information onto Application Insights. You have to enable logging. You also have to ensure that log messages can be correlated to events tracked by Application Insights.

Partial code snippets are given below. You have to complete the missing slots for the code.

```
1     public void ConfigureServices(IServiceCollection services)
2     {
3         services
4             .AddOptions< Slot1 >()
5             .Configure(o => Slot2 = true);
6
7         services.AddMvc();
8     }
9
10
11    public void Configure(IApplicationBuilder app, IHostingEnvironment env, ILoggerFactory loggerFactory)
12    {
13
14        loggerFactory.AddApplicationInsights( Slot3 , LogLevel.Critical);
15
16        app.UseMvc();
17    }
18
19
20
```

Which of the following would go into Slot1?

- A. **IncludeEventId**
- B. **ServerFeatures**
- C. **LoggerFilterOptions**
- D. **ApplicationServices**
- E. **ApplicationInsightsLoggerOptions** ✓
- F. **TrackExceptionsAsExceptionTelemetry**

---

#### Explanation:

Answer – E

This is given as an example in a github article

# Include EventId in logs

It is possible to include `EventId` in telemetry properties. Simply setup `ApplicationInsightsLoggerOptions` instance in `startup.ConfigureServices` method.

```
services  
    .AddOptions<ApplicationInsightsLoggerOptions>()  
    .Configure(o => o.IncludeEventId = true);
```

Since the implementation is clearly given, all other options are incorrect

For more information on logging in ASP.NET Core, one can go to the below link

<https://github.com/Microsoft/ApplicationInsights-aspnetcore/wiki/Logging>

Ask our Experts

Rate this Question?  

Question 34

Correct

Domain :Monitor, troubleshoot, and optimize Azure solutions

A team is developing an ASP.NET Core application. The Application needs to log information onto Application Insights. You have to enable logging. You also have to ensure that log messages can be correlated to events tracked by Application Insights.

Partial code snippets are given below. You have to complete the missing slots for the code.

```
1     public void ConfigureServices(IServiceCollection services)
2     {
3         services
4             .AddOptions< Slot1 >()
5             .Configure(o => Slot2 = true);
6
7         services.AddMvc();
8     }
9
10
11    public void Configure(IApplicationBuilder app, IHostingEnvironment env, ILoggerFactory loggerFactory)
12    {
13
14        loggerFactory.AddApplicationInsights( Slot3 , LogLevel.Critical);
15
16        app.UseMvc();
17    }
18
19
20
```

Which of the following would go into Slot2?

- ✓ A. **IncludeEventId** ✓
- B. **ServerFeatures**
- C. **LoggerFilterOptions**
- D. **ApplicationServices**
- E. **ApplicationInsightsLoggerOptions**
- F. **TrackExceptionsAsExceptionTelemetry**

---

#### Explanation:

Answer - A

This is given as an example in a github article

# Include EventId in logs

It is possible to include `EventId` in telemetry properties. Simply setup `ApplicationInsightsLoggerOptions` instance in `startup.ConfigureServices` method.

```
services
    .AddOptions<ApplicationInsightsLoggerOptions>()
    .Configure(o => o.IncludeEventId = true);
```

Since the implementation is clearly given, all other options are incorrect

For more information on logging in ASP.NET Core, one can go to the below link

<https://github.com/Microsoft/ApplicationInsights-aspnetcore/wiki/Logging>

Ask our Experts

Rate this Question?

Question 35

Correct

Domain :Monitor, troubleshoot, and optimize Azure solutions

A team is developing an ASP.NET Core application. The Application needs to log information onto Application Insights. You have to enable logging. You also have to ensure that log messages can be correlated to events tracked by Application Insights.

Partial code snippets are given below. You have to complete the missing slots for the code.

```
1     public void ConfigureServices(IServiceCollection services)
2     {
3         services
4             .AddOptions< Slot1 >()
5             .Configure(o => Slot2 = true);
6
7         services.AddMvc();
8     }
9
10
11    public void Configure(IApplicationBuilder app, IHostingEnvironment env, ILoggerFactory loggerFactory)
12    {
13
14        loggerFactory.AddApplicationInsights( Slot3 , LogLevel.Critical);
15
16        app.UseMvc();
17    }
18
19
20
```

Which of the following would go into Slot3?

- A. `app.IncludeEventId`
- B. `app.ServerFeatures`
- C. `app.LoggerFilterOptions`
- D. `app.ApplicationServices` ✓
- E. `app.ApplicationInsightsLoggerOptions`
- F. `app.TrackExceptionsAsExceptionTelemetry`

---

#### Explanation:

Answer – D

This is given as an example in a github article

Application Insights SDK for Asp.Net Core provides an extension method `AddApplicationInsights` on `ILoggerFactory` to configure logging. Modify the `startup.cs` class of your application as follows to enable Application Insights logging.

```
public void Configure(  
    IApplicationBuilder app,  
    IHostingEnvironment env,  
    ILoggerFactory loggerFactory)  
{  
    /*...existing code..*/  
    loggerFactory.AddApplicationInsights(app.ApplicationServices, LogLevel.Warning);  
}
```

Since the implementation is clearly given, all other options are incorrect

For more information on logging in ASP.Net Core, one can go to the below link

<https://github.com/Microsoft/ApplicationInsights-aspnetcore/wiki/Logging>

Ask our Experts

Rate this Question?

Question 36

Correct

Domain :Connect to and Consume Azure Services and Third Party Services

Your team has developed an application API based on the OpenAPI specification. You have to ensure that the API can be accessed via an Azure API management service instance. Which of the following Azure PowerShell command would you run to create an API management service?

- A. `Import-AzApiManagementApi -Context $whizlabsApiMgmtContext -SpecificationFormat "Swagger" -SpecificationPath $whizlabsSwaggerPath -Path $whizlabsPath`
- B. `New-AzApiManagementBackend -Context $whizlabsApiMgmtContext -Url $whizlabsurl -Protocol http`
- C. `New-AzApiManagement -ResourceGroupName $whizlabs-rg -Name $whizlabsname -Location $Location -Organization "Whizlabs" -AdminEmail $whizlabsadmin`

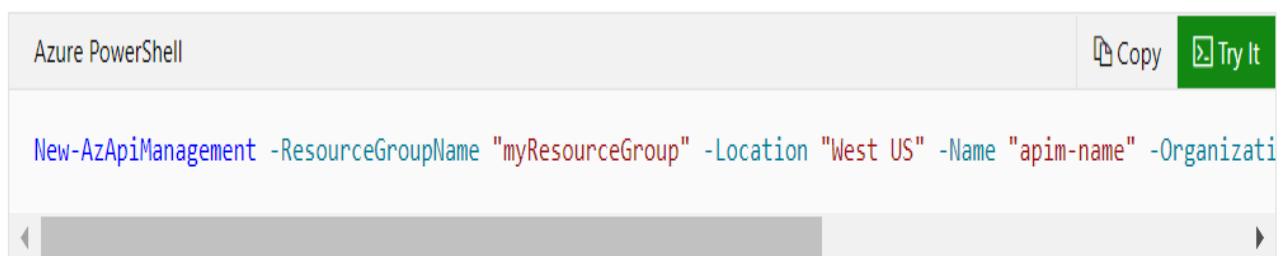
**D. New-AzApiManagementBackendProxy -Url \$whizlabsurl****Explanation:**

Answer – C

First you need to create a new API management instance as shown below

## Create an API Management service

This is a long running operation and could take up to 15 minutes.



Azure PowerShell window showing the command:

```
New-AzApiManagement -ResourceGroupName "myResourceGroup" -Location "West US" -Name "apim-name" -OrganizationId "12345678-1234-1234-1234-1234567890ab"
```

The window includes standard PowerShell UI elements like a scroll bar and status indicators.

Option A is incorrect since this is used to import an Azure API Management API from a file or a URL in Web Application Description Language (WADL), Web Services Description Language (WSDL), or Swagger format

Option B is incorrect since this is used to create a new backend for the API

Option D is incorrect since this is used is just used to create a new backend proxy

For more information on creating an API Instance, one can go to the below link

<https://docs.microsoft.com/en-us/azure/api-management/powershell-create-service-instance>

Ask our Experts

Rate this Question?  

Question 37

Correct

Domain :Develop for Azure Storage

A company is developing an application. The application will be storing data about game scores for players. A class called PlayerScore is in place in the code as a Table Entity. The table is populated with thousands of records.

You need to design code that will retrieve 10 records where the score exceeds 4000.

The following snippets of code have been put in place

```
class PlayerScore : TableEntity
{
    public PlayerScore()
    {
    }

    public PlayerScore(string p_GameID, string p_PlayerID, int p_score, long p_timeplayed)
    {
        this.PartitionKey = p_GameID;
        this.RowKey = p_PlayerID;
        this.score = p_score;
        this.Timeplayed = p_timeplayed;
    }

    public int score { get; set; }
    public long Timeplayed { get; set; }
}

private static void Query()
{
    CloudStorageAccount whizlabs_storage = CloudStorageAccount.Parse(conn_string);
    CloudTableClient whizlabs_table_client = whizlabs_storage.CreateCloudTableClient();
    CloudTable whizlabs_table = whizlabs_table_client.GetTableReference("Player");

    TableQuery<DynamicTableEntity> query = new TableQuery<DynamicTableEntity>().Select
        (new string[] { "score" }).Where(TableQuery.GenerateFilterConditionForInt("score",
        QueryComparisons.GreaterThanOrEqual, 4000)).Take(10);
    EntityResolver<KeyValuePair<string, int?>> resolver = (partitionKey, rowKey, ts, props, etag)
        => new KeyValuePair<string, int?>(rowKey, props["score"].Int32Value);

    foreach (var scoreItem in whizlabs_table.ExecuteQuery(query, resolver, null, null))
    {
        Console.WriteLine(scoreItem.Key);
        Console.WriteLine(scoreItem.Value);
    }
}
```

Does the code query the Azure Table and retrieve the TimePlayed property from the table?

A. Yes

✓ B. No ✓

**Explanation:**

Answer - B

No, since the Select part is in place, this will only select the score property.

```
TableQuery<DynamicTableEntity> query = new TableQuery<DynamicTableEntity>().  
    Select(new string[] { "score" }).Where(TableQuery.GenerateFilterConditionForInt("score",  
        QueryComparisons.GreaterThanOrEqual, 4000)).Take(10);
```

For more information on table operations in .Net, one can go to the below link

<https://docs.microsoft.com/en-us/azure/cosmos-db/table-storage-how-to-use-dotnet>

Ask our Experts

Rate this Question?  

Question 38

Correct

Domain :Develop for Azure Storage

A company is developing an application. The application will be storing data about game scores for players. A class called PlayerScore is in place in the code as a Table Entity. The table is populated with thousands of records.

You need to design code that will retrieve 10 records where the score exceeds 4000.

The following snippets of code have been put in place

```
class PlayerScore : TableEntity
{
    public PlayerScore()
    {
    }
    public PlayerScore(string p_GameID, string p_PlayerID, int p_score, long p_timeplayed)
    {
        this.PartitionKey = p_GameID;
        this.RowKey = p_PlayerID;
        this.score = p_score;
        this.Timeplayed = p_timeplayed;
    }
    public int score { get; set; }
    public long Timeplayed { get; set; }
}
```

```
private static void Query()
{
    CloudStorageAccount whizlabs_storage = CloudStorageAccount.Parse(conn_string);
    CloudTableClient whizlabs_table_client = whizlabs_storage.CreateCloudTableClient();
    CloudTable whizlabs_table = whizlabs_table_client.GetTableReference("Player");

    TableQuery<DynamicTableEntity> query = new TableQuery<DynamicTableEntity>().Select
    (new string[] { "score" }).Where(TableQuery.GenerateFilterConditionForInt("score",
    QueryComparisons.GreaterThanOrEqual, 4000)).Take(10);
    EntityResolver<KeyValuePair<string, int?>> resolver = (partitionKey, rowKey, ts, props, etag)
=> new KeyValuePair<string, int?>(rowKey, props["score"].Int32Value);

    foreach (var scoreItem in whizlabs_table.ExecuteQuery(query, resolver, null, null))
    {
        Console.WriteLine(scoreItem.Key);
        Console.WriteLine(scoreItem.Value);
    }
}
```

Does the code return a maximum of ten records?

- A. Yes
- B. No

**Explanation:**

Answer – A

Yes, because of the Take condition which limits the rows, the number of records will be limited to 10.

```
TableQuery<DynamicTableEntity> query = new TableQuery<DynamicTableEntity>().  
Select(new string[] { "score" }).Where(TableQuery.GenerateFilterCondition("score",  
QueryComparisons.GreaterThanOrEqual, 4000)).Take(10);|
```



For more information on table operations in .Net, one can go to the below link

<https://docs.microsoft.com/en-us/azure/cosmos-db/table-storage-how-to-use-dotnet>

Ask our Experts

Rate this Question?

Question 39

Correct

Domain :Develop for Azure Storage

A company is developing an application. The application will be storing data about game scores for players. A class called PlayerScore is in place in the code as a Table Entity. The table is populated with thousands of records.

You need to design code that will retrieve 10 records where the score exceeds 4000.

The following snippets of code have been put in place

```
class PlayerScore : TableEntity
{
    public PlayerScore()
    {
    }

    public PlayerScore(string p_GameID, string p_PlayerID, int p_score, long p_timeplayed)
    {
        this.PartitionKey = p_GameID;
        this.RowKey = p_PlayerID;
        this.score = p_score;
        this.Timeplayed = p_timeplayed;
    }

    public int score { get; set; }
    public long Timeplayed { get; set; }
}
```

```
private static void Query()
{
    CloudStorageAccount whizlabs_storage = CloudStorageAccount.Parse(conn_string);
    CloudTableClient whizlabs_table_client = whizlabs_storage.CreateCloudTableClient();
    CloudTable whizlabs_table = whizlabs_table_client.GetTableReference("Player");

    TableQuery<DynamicTableEntity> query = new TableQuery<DynamicTableEntity>().Select
    (new string[] { "score" }).Where(TableQuery.GenerateFilterConditionForInt("score",
    QueryComparisons.GreaterThanOrEqual, 4000)).Take(10);
    EntityResolver<KeyValuePair<string, int?>> resolver = (partitionKey, rowKey, ts, props, etag)
=> new KeyValuePair<string, int?>(rowKey, props["score"].Int32Value);

    foreach (var scoreItem in whizlabs_table.ExecuteQuery(query, resolver, null, null))
    {
        Console.WriteLine(scoreItem.Key);
        Console.WriteLine(scoreItem.Value);
    }
}
```

Does the code return all records to the client? The client will then display the records where the score is greater than 4000?

- A. Yes
- B. No

**Explanation:**

Answer – A

Here since the query is performed on a property that is not related to the Partition Key, all the rows from the table will be fetched. The Microsoft documentation mentions the following

- A **Table Scan** does not include the **PartitionKey** and is very inefficient because it searches all of the partitions that make up your table in turn for any matching entities. It will perform a table scan regardless of whether or not your filter uses the **RowKey**. For example: `$filter=LastName eq 'Jones'`

For more information on designing table storage, one can go to the below link

<https://docs.microsoft.com/en-us/azure/storage/tables/table-storage-design-for-query>

Ask our Experts

Rate this Question?  

Question 40

Correct

Domain :Connect to and Consume Azure Services and Third Party Services

[View Case Study](#)

You have to ensure that all sign-in and sign-out events can be processed by the EventGridController.

You propose the following solution

**"Create separate Azure Event Grid topics and subscriptions for sign-in and sign-out events."**

Does this solution meet the requirement?

- ✓ A. Yes 
- B. No

**Explanation:**

Answer - A

Yes, you can create topics for both the sign in and sign out events.

For more information on posting to a custom topic, one can go to the below link

<https://docs.microsoft.com/en-us/azure/event-grid/post-to-custom-topic>

[Ask our Experts](#)Rate this Question?  [Question 41](#)

Correct

**Domain :Connect to and Consume Azure Services and Third Party Services**[View Case Study](#)

You have to ensure that all sign-in and sign-out events can be processed by the EventGridController.

You propose the following solution

**"Create a new Azure Event Grid subscription for all authentication that delivers messages to an Azure Event Hub. Use the subscription to process sign-out events"**

Does this solution meet the requirement?

A. Yes

 B. No 

**Explanation:**

Answer – B

The Azure Event Hub is normally used as a data ingestion service

For more information on Event Hubs, one can go to the below link

<https://docs.microsoft.com/en-us/azure/event-hubs/event-hubs-about>

[Ask our Experts](#)Rate this Question?  [Question 42](#)

Correct

**Domain :Connect to and Consume Azure Services and Third Party Services**[View Case Study](#)

You have to ensure that all sign-in and sign-out events can be processed by the EventGridController.

You propose the following solution

**"Create a new Azure Event Grid topic and add a subscription for the events."**

Does this solution meet the requirement?

A. Yes

✓ B. No 

---

### Explanation:

Answer – B

It is better to create separate topics. Since the sign-outs need to be processed immediately, you should create a separate topic for the sign in a separate topic for the sign-out process.

For more information on posting to a custom topic, one can go to the below link

<https://docs.microsoft.com/en-us/azure/event-grid/post-to-custom-topic>

---

Ask our Experts

Rate this Question?  

Question 43

Correct

Domain :Connect to and Consume Azure Services and Third Party Services

[View Case Study](#)

You are designing the class that will be used to parse the Event Data from the Event Grid. You have to complete the below class segment

```
public class LoginEvent
{
    public string Slot1 { get; set; }

    public string Slot2 { get; set; }

    public string Slot3 { get; set; }

    public string subject { get; set; }
    public DateTime eventTime { get; set; }
    public Dictionary<string, string> data { get; set; }
    public string Serialized()
    {
        return JsonConvert.SerializeObject(this);
    }
}
```

Which of the following will go into Slot1?

- ✓ A. id
- B. eventType
- C. topic
- D. metadataVersion

#### Explanation:

Answer – A

Since the EventGridController.cs file refers to the id and topic value, we need to ensure these are in place in the class definition for the object

The Microsoft documentation below specifies the Event Schema for the Event Grid

N

## Event schema

The following example shows the properties that are used by all event publishers:

```
JSON Copy  
[  
  {  
    "topic": string,  
    "subject": string,  
    "id": string,  
    "eventType": string,  
    "eventTime": string,  
    "data":{  
      object-unique-to-each-publisher  
    },  
    "dataVersion": string,  
    "metadataVersion": string  
  }  
]
```

For more information on the Event Grid schema, one can go to the below link

<https://docs.microsoft.com/en-us/azure/event-grid/event-schema>

Ask our Experts

Rate this Question?  

Question 44

Correct

Domain :Connect to and Consume Azure Services and Third Party Services

[View Case Study](#)

You are designing the class that will be used to parse the Event Data from the Event Grid. You have to complete the below class segment

```
public class LoginEvent
{
    public string Slot1 { get; set; }

    public string Slot2 { get; set; }

    public string Slot3 { get; set; }

    public string subject { get; set; }
    public DateTime eventTime { get; set; }
    public Dictionary<string, string> data { get; set; }
    public string Serialized()
    {
        return JsonConvert.SerializeObject(this);
    }
}
```

Which of the following will go into Slot2?

- A. id
- B. eventType
- C. topic
- D. metadataVersion

#### Explanation:

Answer – C

Since the EventGridController.cs file refers to the id and topic value, we need to ensure these are in place in the class definition for the object

The Microsoft documentation below specifies the Event Schema for the Event Grid

N

## Event schema

The following example shows the properties that are used by all event publishers:

```
JSON Copy  
[  
  {  
    "topic": string,  
    "subject": string,  
    "id": string,  
    "eventType": string,  
    "eventTime": string,  
    "data":{  
      object-unique-to-each-publisher  
    },  
    "dataVersion": string,  
    "metadataVersion": string  
  }  
]
```

For more information on the Event Grid schema, one can go to the below link

<https://docs.microsoft.com/en-us/azure/event-grid/event-schema>

Ask our Experts

Rate this Question?  

Question 45

Correct

Domain :Connect to and Consume Azure Services and Third Party Services

[View Case Study](#)

You are designing a class that will be used to parse the Event Data from the Event Grid. You have to complete the below class segment

```
public class LoginEvent
{
    public string Slot1 { get; set; }

    public string Slot2 { get; set; }

    public string Slot3 { get; set; }

    public string subject { get; set; }
    public DateTime eventTime { get; set; }
    public Dictionary<string, string> data { get; set; }
    public string Serialized()
    {
        return JsonConvert.SerializeObject(this);
    }
}
```

Which of the following will go into Slot3?

- A. id
- B. eventType
- C. topic
- D. metadataVersion

#### Explanation:

Answer – B

We also need to have the eventType in place.

The Microsoft documentation below specifies the Event Schema for the Event Grid

N

## Event schema

The following example shows the properties that are used by all event publishers:

```
JSON Copy  
[  
  {  
    "topic": string,  
    "subject": string,  
    "id": string,  
    "eventType": string,  
    "eventTime": string,  
    "data":{  
      object-unique-to-each-publisher  
    },  
    "dataVersion": string,  
    "metadataVersion": string  
  }  
]
```

For more information on the Event Grid schema, one can go to the below link

<https://docs.microsoft.com/en-us/azure/event-grid/event-schema>

---

Ask our Experts

Rate this Question?

Question 46

Correct

Domain :Develop Azure Platform as a Service Compute Solutions

[View Case Study](#)

Users report that anomaly detection emails can sometimes arrive several minutes after an anomaly is detected. How can you resolve this issue?

- A. Ensure that the Azure Function is using an App Service plan.
- B. Set Always On to false

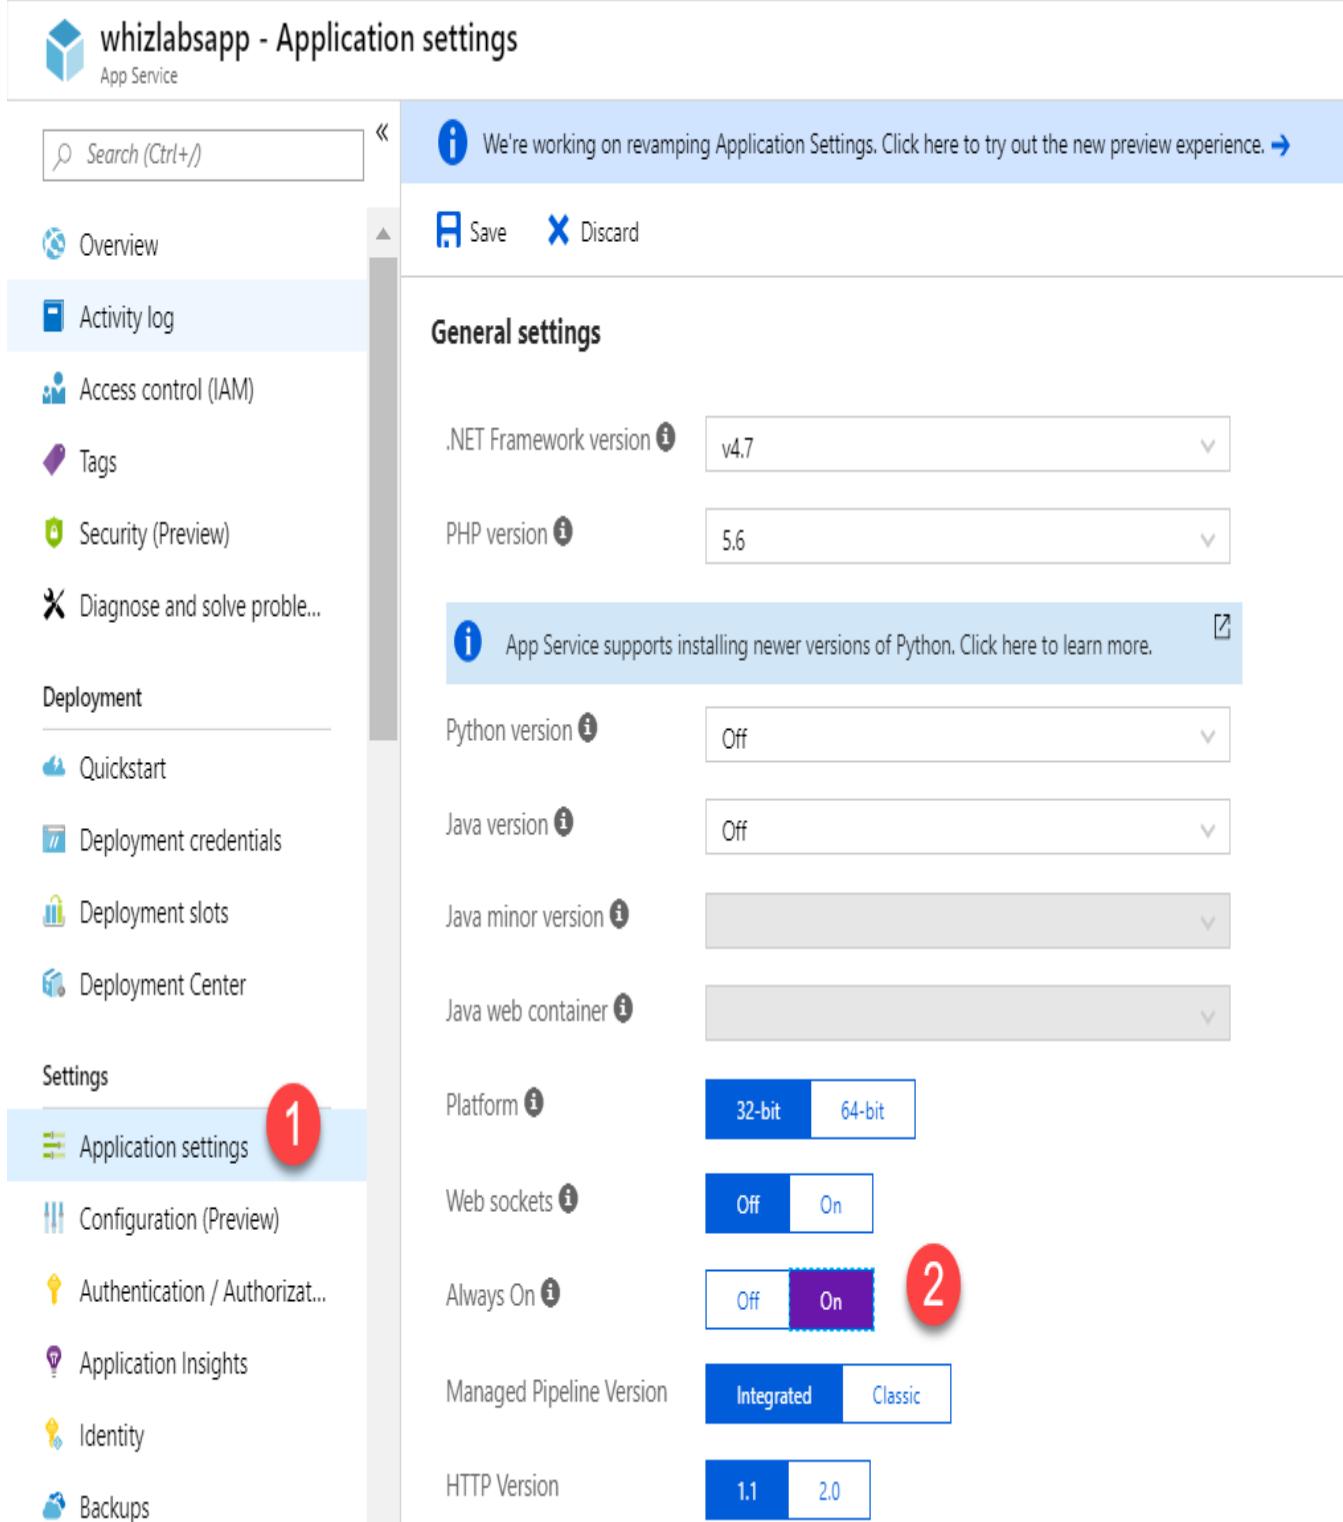
C. Ensure that the Azure Function is set to use a consumption plan.

- ✓ D. Set Always On to true. 

**Explanation:**

Answer – D

Here the issue is that the Azure Web app is being stopped when it is not being used. For this you have to ensure the AlwaysOn setting for the Web App is implemented as True as shown below



The screenshot shows the Azure Application Settings page for the 'whizlabsapp' App Service. The left sidebar lists various settings categories, with 'Application settings' selected and highlighted by a red circle labeled '1'. The main area displays 'General settings' with the following configuration:

- .NET Framework version: v4.7
- PHP version: 5.6
- Python version: Off
- Java version: Off
- Java minor version: (dropdown menu)
- Java web container: (dropdown menu)
- Platform: 32-bit (selected)
- Web sockets: On
- Always On: On (highlighted by a red circle labeled '2')
- Managed Pipeline Version: Integrated
- HTTP Version: 1.1

The Microsoft documentation mentions the following on the setting.

**Always On.** By default, apps are unloaded if they are idle for some period of time. This lets the system conserve resources. In Basic or Standard mode, you can enable **Always On** to keep the app loaded all the time. If your app runs continuous WebJobs or runs WebJobs triggered using a CRON expression, you should enable **Always On**, or the web jobs may not run reliably.

Options A and C are incorrect since this is not an issue with the Azure Function, but with the Web App

Option B is incorrect since the setting should be True

For more information on Azure Web App settings, one can go to the below link

<https://docs.microsoft.com/en-us/azure/app-service/web-sites-configure>

---

Ask our Experts

Rate this Question?  

Question 47

Correct

Domain :Monitor, troubleshoot, and optimize Azure solutions

[View Case Study](#)

The "WhizlabsPolicyService" application must be able to scale on demand. Which Azure Application Insights data model should you use?

- ✓ A. An Application Insights metric 
- B. An Application Insights dependency
- C. An Application Insights trace
- D. An Application Insights event

---

**Explanation:**

Answer – A

You can use Application Insights metrics to scale Web Apps. The Microsoft documentation gives an example on this where the metrics source is Application Insights.

The screenshot shows the Azure portal interface for managing an autoscale setting. On the left, there's a sidebar with various icons. The main area displays the 'Autoscale setting' for 'Web api autoscale' within a resource group 'contoso-web'. It shows an instance count of 1. Below this, there's a 'Default' scale condition with a note about scaling based on a metric (CPU percentage above 70%). There are buttons for '+ Add a rule' and '+ Add a scale condition'. On the right, a modal window titled 'Scale rule' is open, detailing the configuration for this specific rule. The 'Metric source' is set to 'Application Insights', and the 'Resource' is 'loans-app-ai'. The 'Criteria' section is expanded, showing the following settings:

- \* Time aggregation: Total
- \* Metric name: LoanSubmissions
- \* Time grain (in mins): 1
- \* Operator: Greater than
- \* Threshold: 100
- \* Duration (in minutes): 10

Since this is clearly mentioned in the Microsoft documentation, all other options are incorrect

For more information on autoscaling based on a custom metric, one can go to the below link

<https://docs.microsoft.com/en-us/azure/azure-monitor/platform/autoscale-custom-metric>

#### Ask our Experts

Rate this Question?

#### Question 48

Correct

#### Domain :Connect to and Consume Azure Services and Third Party Services

##### [View Case Study](#)

There is a plan to use Azure Redis Cache to improve the performance of the "WhizlabsPolicyService" application. Which of the following would you store in Azure Redis Cache?

- A. **HttpContext.Items**
- B. **ViewState**

✓ C. Session state 

D. TempData

---

**Explanation:**

Answer – C

You would ideally store the session state in Azure Redis. The Microsoft documentation mentions the following as one of the patterns or use cases for using Azure Redis.

User session caching This pattern is commonly used with shopping carts and other user history type information that a web application may want to associate with user cookies. Storing too much in a cookie can have a negative impact on performance as the cookie size grows and is passed and validated with every request. A typical solution is to use the cookie as a key to query the data in a backend database. Using an in-memory cache, like Azure Cache for Redis, to associate information with a user is much faster than interacting with a full relational database.

For more information on Azure Redis, one can go to the below link

<https://docs.microsoft.com/en-us/azure/azure-cache-for-redis/cache-overview>

Ask our Experts

Rate this Question?  

Question 49

Correct

Domain :Implement Azure Security

A company maintains an existing Azure SQL Database. Keys which can be used to encrypt the database is stored in the Azure Key Vault. The database contains the following columns

FirstName

LastName

Age

SSN

An external consulting firm has been given the responsibility to manage the database. But you have to ensure that the external party does not have the ability to access the data in the SSN column of the table

You decide on using the following protection method

**"Enable AlwaysOn encryption"**

Would this protection method fulfil the requirement?

✓ A. Yes 

B. No

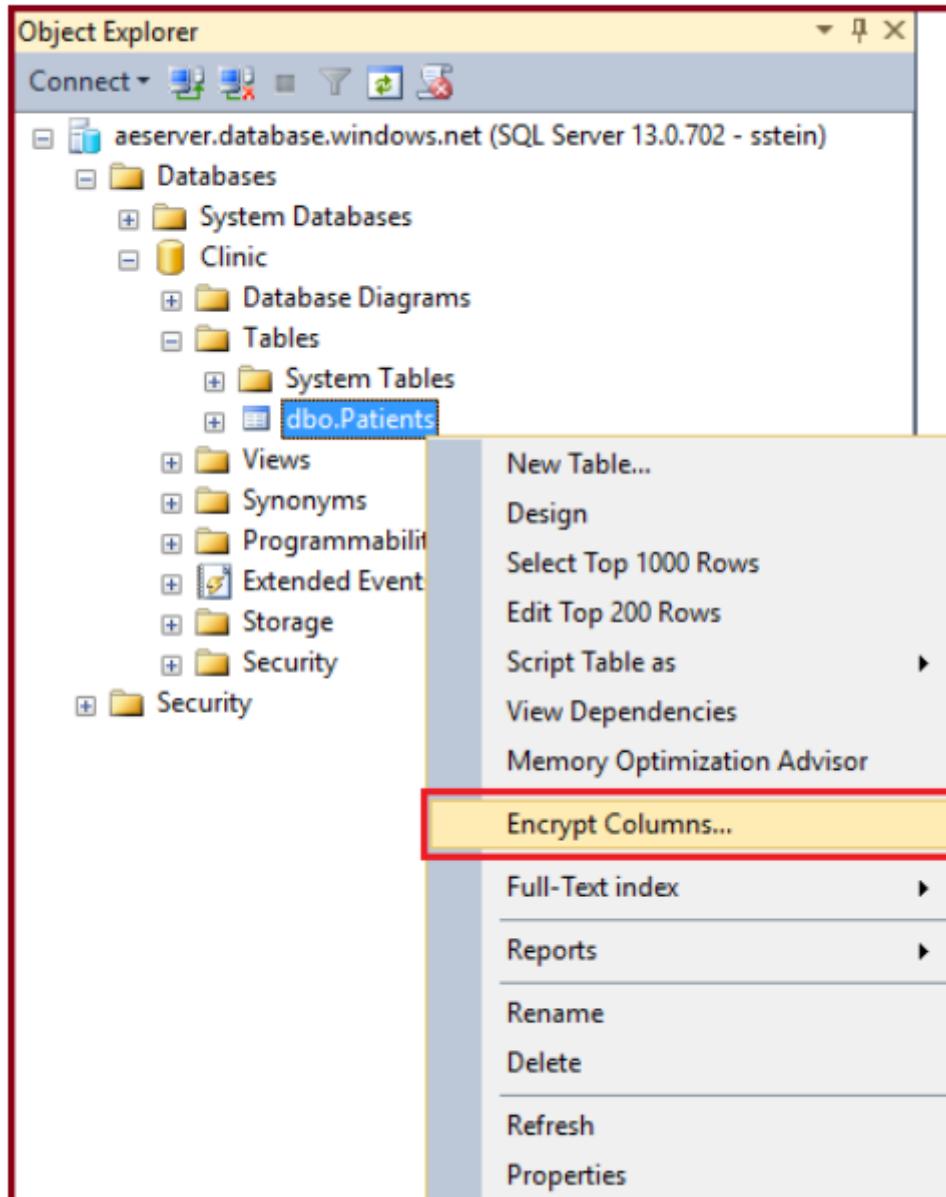
**Explanation:**

Answer – A

A similar example of this is given in the Microsoft documentation as shown below. So, you should ensure AlwaysOn Encryption for this requirement

## Encrypt columns (configure Always Encrypted)

SSMS provides a wizard that helps you easily configure Always Encrypted by setting up the column master key, column encryption key, and encrypted columns for you.



**Always Encrypted**

**Column Selection**

**Introduction** **Column Selection** **Master Key Configuration** **Validation** **Summary** **Results** **Help**

Search column name...

Apply one key to all checked columns: **CEK\_Auto1 (New)**

Name	State	Encryption Type	Encryption Key
dbo.Patients			
PatientId			
<input checked="" type="checkbox"/> SSN	Deterministic	CEK_Auto1 (New)	<b>Choose Type...</b>
<input type="checkbox"/> FirstName	Deterministic		
<input type="checkbox"/> LastName			
<input type="checkbox"/> MiddleName			
StreetAddress			
City			
ZipCode			
State			
<input checked="" type="checkbox"/> BirthDate	Randomized	CEK_Auto1 (New)	

For more information on AlwaysOn Encryption for Azure SQL Database, please go ahead and visit the below URL

<https://docs.microsoft.com/en-us/azure/sql-database/sql-database-always-encrypted-azure-key-vault>

Ask our Experts

Rate this Question?

Question 50

Correct

Domain :Implement Azure Security

A company maintains an existing Azure SQL Database. Keys which can be used to encrypt the database is stored in the Azure Key Vault. The database contains the following columns

FirstName

LastName

.Age

SSN

An external consulting firm has been given the responsibility to manage the database. But you have to ensure that the external party does not have the ability to access the data in the SSN column of the table

You decide on using the following protection method

**"Set the column encryption setting as disabled"**

Would this protection method fulfil the requirement?

A. Yes

✓ B. No 

---

### Explanation:

Answer - B

You have to setup AlwaysOn Encryption. This is how you can protect certain columns in the Azure SQL Database

For more information on AlwaysOn Encryption for Azure SQL Database, please go ahead and visit the below URL

<https://docs.microsoft.com/en-us/azure/sql-database/sql-database-always-encrypted-azure-key-vault>

---

Ask our Experts

Rate this Question?  

Question 51

Correct

Domain :Implement Azure Security

A company maintains an existing Azure SQL Database. Keys which can be used to encrypt the database is stored in the Azure Key Vault. The database contains the following columns

FirstName

LastName

.Age

SSN

An external consulting firm has been given the responsibility to manage the database. But you have to ensure that the external party does not have the ability to access the data in the SSN column of the table

You decide on using the following protection method

**"Assign users to the Public fixed database role"**

Would this protection method fulfil the requirement?

A. Yes

✓ B. No 

---

### Explanation:

Answer - B

Roles are more from a permission level aspect from the entire database perspective and not from a column to column perspective

For more information on security for SQL databases, please go ahead and visit the below URL

<https://docs.microsoft.com/en-us/sql/relational-databases/security/authentication-access/getting-started-with-database-engine-permissions?view=sql-server-2017>

Ask our Experts

Rate this Question?  

---

### Question 52

Correct

Domain :Implement Azure Security

A company maintains an existing Azure SQL Database. Keys which can be used to encrypt the database is stored in the Azure Key Vault. The database contains the following columns

FirstName

LastName

Age

SSN

An external consulting firm has been given the responsibility to manage the database. But you have to ensure that the external party does not have the ability to access the data in the SSN column of the table

You decide on using the following protection method

**"Store column encryption keys in the system catalogue view of the database"**

Would this protection method fulfil the requirement?

A. Yes

✓ B. No 

---

### Explanation:

Answer – B

The keys should always be stored in the Azure Key vault from a security perspective and should not be stored in the database itself.

The Microsoft documentation mentions the following

## Create a key vault to store your keys

Now that your client app is configured and you have your application ID, it's time to create a key vault and configure its access policy so you and your application can access the vault's secrets (the Always Encrypted keys). The *create, get, list, sign, verify, wrapKey, and unwrapKey* permissions are required for creating a new column master key and for setting up encryption with SQL Server Management Studio.

For more information on AlwaysOn Encryption for Azure SQL Database, please go ahead and visit the below URL

<https://docs.microsoft.com/en-us/azure/sql-database/sql-database-always-encrypted-azure-key-vault>

Ask our Experts

Rate this Question?  

Question 53

Correct

Domain :Develop Azure Platform as a Service Compute Solutions

A development team is developing an application that needs to be deployed to an Azure Virtual Machine. It needs to be ensured that the underlying disk is encrypted using a Key from the Azure Key Vault service. Which of the following powershell cmdlets would you execute for this requirement? Choose 3 answers from the options given below

- ✓ A. New-AzVM 
- B. New-AzVmDiskEncryption
- ✓ C. Set-AzVmDiskEncryptionExtension 
- D. Set-AzVM
- ✓ E. Get-AzKeyVault 

### Explanation:

Answer – A, C and E

First ensure that the Virtual Machine is in place.

Also ensure that you have an Azure Key vault already in place.

And then finally use the Set-AzVmDiskEncryptionExtension to enable encryption on the disk

An example of this is given in the Microsoft documentation

- **Encrypt a running VM:** The script below initializes your variables and runs the Set-AzVmDiskEncryptionExtension cmdlet. The resource group, VM, and key vault should have already been created as prerequisites. Replace MyKeyVaultResourceGroup, MyVirtualMachineResourceGroup, MySecureVM, and MySecureVault with your values.

Azure PowerShell Copy

```
$KVRGname = 'MyKeyVaultResourceGroup';
$VMRGName = 'MyVirtualMachineResourceGroup';
$vmName = 'MySecureVM';
$keyVaultName = 'MySecureVault';
$keyVault = Get-AzKeyVault -VaultName $keyVaultName -ResourceGroupName $KVRGname;
$diskEncryptionKeyVaultUrl = $keyVault.VaultUri;
$keyVaultResourceId = $keyVault.ResourceId;

Set-AzVmDiskEncryptionExtension -ResourceGroupName $VMRGName -VMName $vmName -DiskEncryptionKeyVault
```

Option B is invalid since this is not a valid cmdlet

Option D is invalid since this is used to set a VM as generalized

For more information on encrypting a VM, please go ahead and visit the below URL

<https://docs.microsoft.com/en-us/azure/security/quick-encrypt-vm-powershell>

Ask our Experts

Rate this Question? 😊 😢

Question 54

Correct

Domain :Develop Azure Platform as a Service Compute Solutions

A consultant needs to deploy Web Applications to the Azure Web App service for 4 customers. Each customer needs to have the application running on a separate individual instance. The

following key requirements are also in place

- .Ability to automatically scale on demand
- .Ability to use deployment slots to test staging environments
- .All Azure resources should be located in a separate isolated network
- Costs need to be minimized

Which of the following would you choose for the App Service Plan?

- A. Basic
- B. Standard
- C. Premium
- ✓ D. Isolated 

### Explanation:

Answer – D

Here the defining requirement is to ensure that Azure resources are located in a separate environment, and for that we need to use the Isolated App Service Plan

The Microsoft documentation mentions the following

#### Isolated Service Plan

The Isolated service plan is designed to run mission critical workloads, that are required to run in a virtual network. The Isolated plan allows customers to run their apps in a private, dedicated environment in an Azure datacenter using Dv2-series VMs with faster processors, SSD storage, and double the memory-to-core ratio compared to Standard. The private environment used with an Isolated plan is called the App Service Environment. The plan can scale to 100 instances with more available upon request. You can find more details on the Isolated plan and [App Service Environments](#). In addition to the price per Isolated plan instance there is also a flat fee for each App Service Environment of \$1.43/hour (~\$1,043.82/month).

INSTANCE	CORES	RAM	STORAGE	PRICES
I1	1	3.50 GB	1 TB	\$0.30/hour
I2	2	7 GB	1 TB	\$0.60/hour
I3	4	14 GB	1 TB	\$1.20/hour

The Isolated App Service plan supports Auto Scaling and Deployment slots

Since this is the only viable option, all other options are incorrect

For more information on App Service Plans, please go ahead and visit the below URL

<https://azure.microsoft.com/en-us/pricing/details/app-service/plans/>

[Ask our Experts](#)Rate this Question?  

Question 55

Correct

**Domain :Develop Azure Platform as a Service Compute Solutions**

A consultant needs to deploy Web Applications to the Azure Web App service for 4 customers. Each customer needs to have the application running on a separate individual instance. The following key requirements are also in place

- .Ability to automatically scale on demand
- .Ability to use deployment slots to test staging environments
- .All Azure resources should be located in a separate isolated network
- Costs need to be minimized

How many instances would you keep running for the requirement?

- ✓ A. 4 
- B. 8
- C. 12
- D. 16

**Explanation:**

Answer – A

Since each we need to ensure each customer's Web app is running on a separate instance and since we also need to cut down on costs , we should limit the number of instances to 4.

Since this is the only viable option, all other options are incorrect

For more information on App Service Plans, please go ahead and visit the below URL

<https://azure.microsoft.com/en-us/pricing/details/app-service/plans/>

[Ask our Experts](#)Rate this Question?  [Finish Review](#)

**Certification**

Cloud Certification  
Java Certification  
PM Certification  
Big Data Certification

**Company**

Support  
Discussions  
Blog

**Follow us**

---

© Copyright 2019. Whizlabs Software Pvt. Ltd. All Right Reserved.